



# Computer Vision: Local transformations

Nicolás Pérez de la Blanca,  
[nicolas@decsai.ugr.es](mailto:nicolas@decsai.ugr.es)

# Plan for today

- Image noise
- Local transformations
  - Examples: smoothing filters
- Convolution / correlation
- Edge detection

# Readings for today

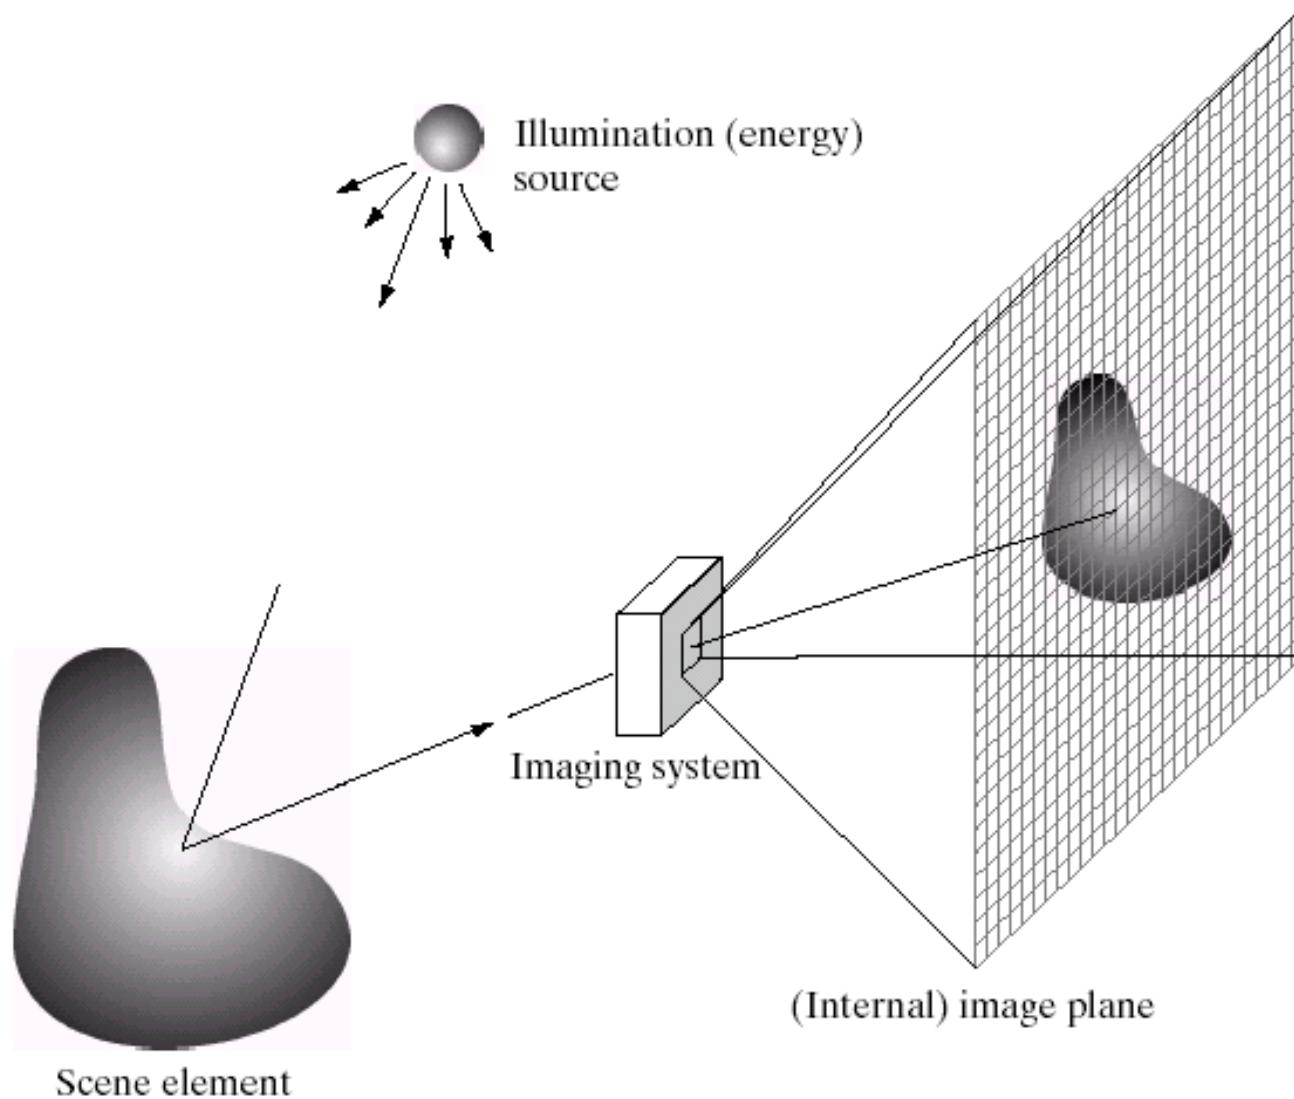
- ***Image Filtering***
  - Readings: Szel., Ch. 3.1-3.3
- ***Edge detection***
  - Readings: Szel., Ch. 4.2

# What is an image?



**Very complex data encoding multiples cues**

# Image Formation



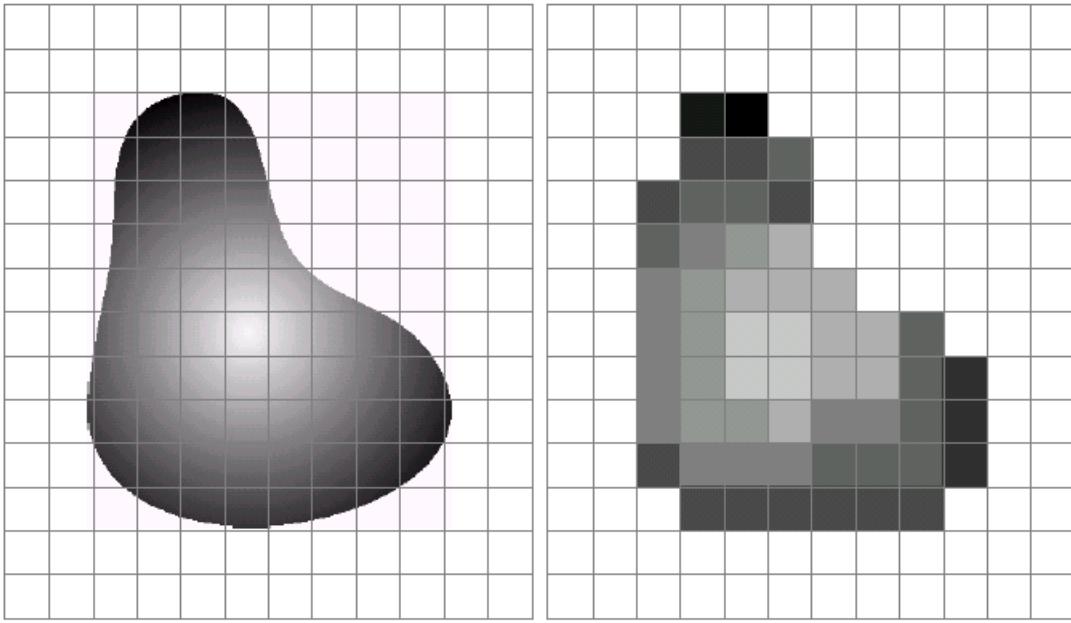
# Digital camera



A digital camera replaces film with a sensor array

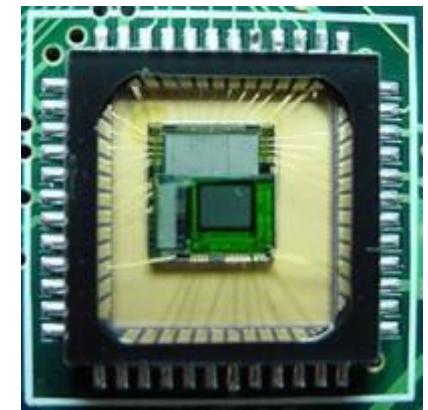
- Each cell in the array is light-sensitive diode that converts photons to electrons
- <http://electronics.howstuffworks.com/digital-camera.htm>

# Digital images



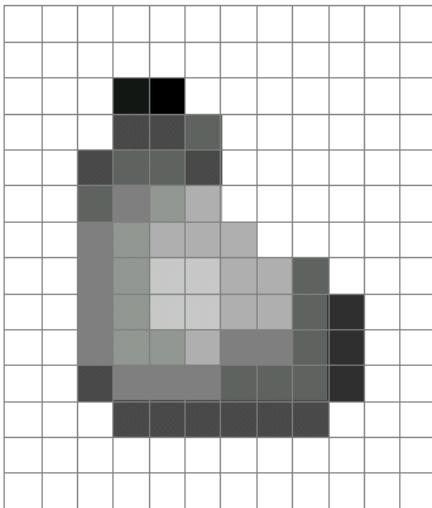
a b

**FIGURE 2.17** (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.



# What is an image?

- A grid (matrix) of intensity values

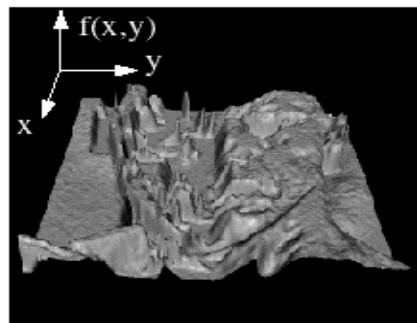


255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	20	0	255	255	255	255	255	255	255	255	255	255
255	255	255	75	75	75	255	255	255	255	255	255	255	255	255	255
255	255	75	95	95	75	255	255	255	255	255	255	255	255	255	255
255	255	96	127	145	175	255	255	255	255	255	255	255	255	255	255
255	255	127	145	175	175	175	255	255	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	95	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	95	47	255	255	255	255	255	255
255	255	127	145	145	175	127	127	95	47	255	255	255	255	255	255
255	255	74	127	127	127	127	95	95	95	47	255	255	255	255	255
255	255	255	74	255	74	74	74	74	74	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255

(common to use one byte per value: 0 = black, 255 = white)

# Digital images

- **Sample** the 2D space on a regular grid
- **Quantize** each sample (round to nearest integer)
- Image thus represented as a matrix of integer values.

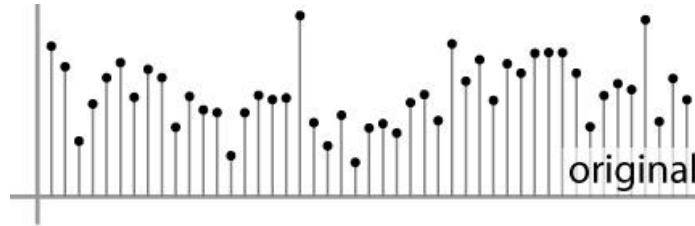


$j$  →

$i$  ↓

62	79	23	119	120	105	4	0
10	10	9	62	12	78	34	0
10	58	197	46	46	0	0	48
176	135	5	188	191	68	0	49
2	1	1	29	26	37	0	77
0	89	144	147	187	102	62	208
255	252	0	166	123	62	0	31
166	63	127	17	1	0	99	30

2D



1D

# Digital color images

Color images,  
RGB color  
space



R

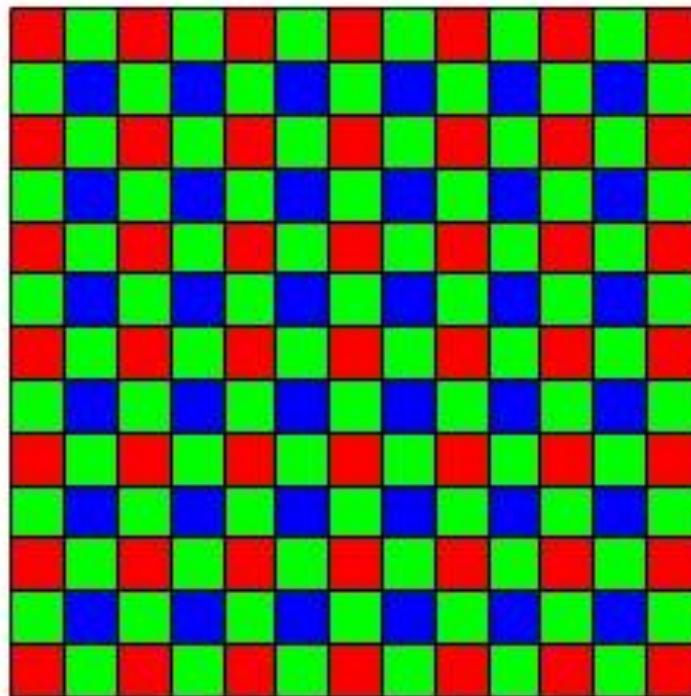


G



B

# Digital color images



**Bayer filter**

© 2000 How Stuff Works

# Images in OpenCV

- Image represented as a matrix `im`
- `im=imread(filename)` returns a CV\_8UC3 image (values 0 to 255)
  - Convert to float format with `im.convertTo()`

row \ column	0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99
0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91	
0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92	
0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95	
0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85	
0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33	
0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74	
0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93	
0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99	
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97	
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93	
	0.95	0.45	0.58	0.88	0.45	0.42	0.77	0.75	0.71	0.90	0.99
	0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
	0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93
	0.95	0.45	0.58	0.88	0.45	0.42	0.77	0.75	0.71	0.90	0.99
	0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
	0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93
	0.95	0.45	0.58	0.88	0.45	0.42	0.77	0.75	0.71	0.90	0.99
	0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
	0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93

# What is the most relevant information cue ?

- Pixel and nearby regions show dependence between their values
  - No rule for distant pixels/regions (high spatial variability)
- Main approach: To characterize images from local informations.

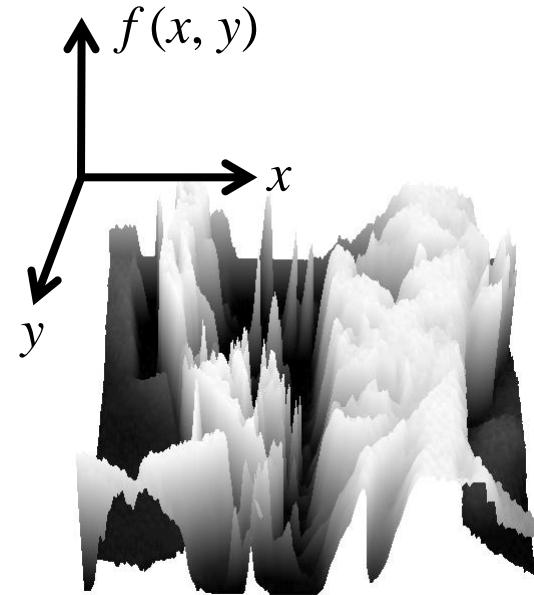


# The image as a function

- We can think of a (grayscale) image as a **function**,  $f$ , from  $\mathbb{R}^2$  to  $\mathbb{R}$ :
  - $f(x, y)$  gives the **intensity** at position  $(x, y)$



[snoop](#)



[3D view](#)

- A **digital** image is a discrete (**sampled, quantized**) version of this function

# Basic Image Transformations

- In order to extract useful information to approach our goals we apply transformations on the image values
- **Local** : Image filtering using local masks
- **Local/Global**: Geometric transformations

# Examples of transformations

- As with any function, we can apply operators to an image



$$g(x,y) = f(x,y) + 20$$



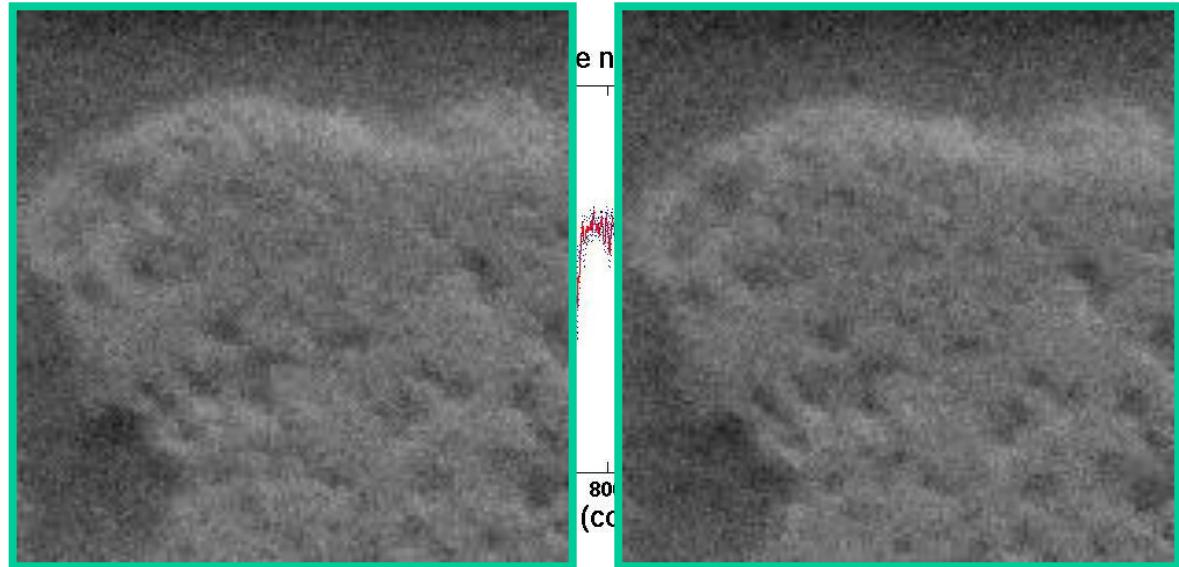
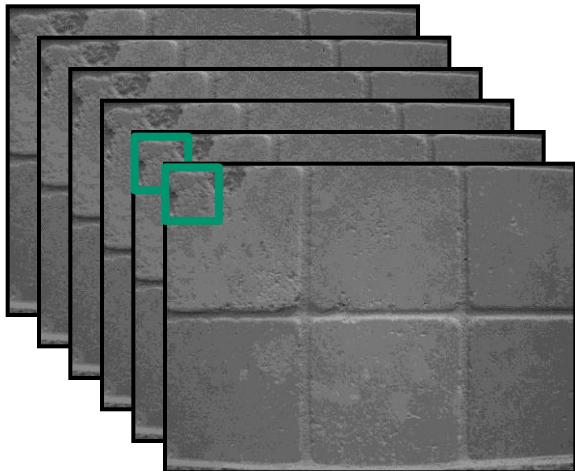
$$g(x,y) = f(-x,y)$$

- We'll talk about a special kind of operator, *convolution* (Local transformation)

# Image filtering

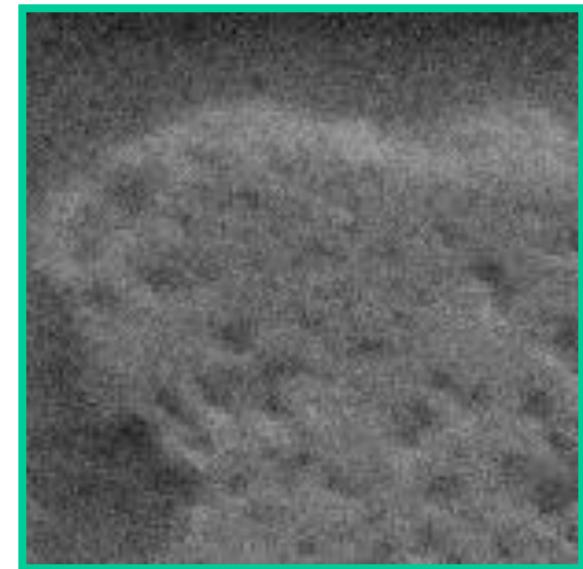
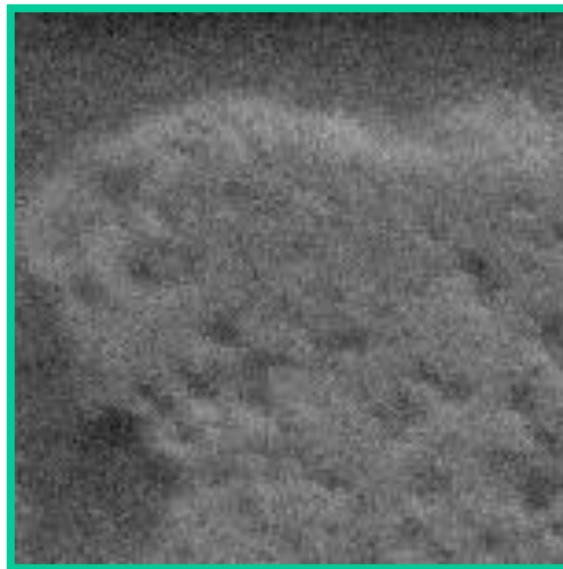
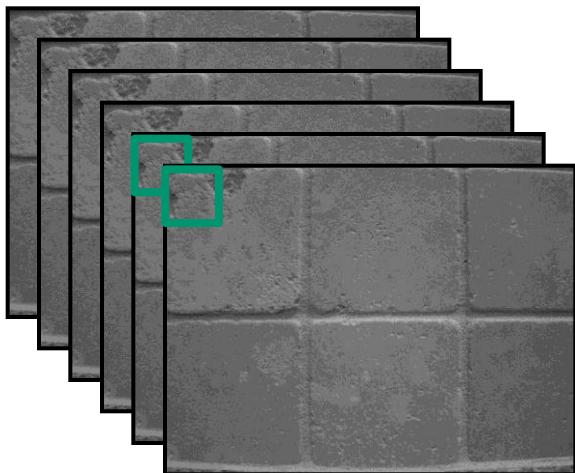
- Compute a function of the local neighborhood at each pixel in the image
  - Function specified by a “filter” or mask saying how to combine values from neighbors.
  - We extract only a part of the image information
- Uses of filtering:
  - Enhance an image (denoise, resize, etc)
  - Extract information (texture, edges, etc)
  - Detect patterns (template matching)

# Simple case: noise reduction



- Even multiple images of the **same static scene** will not be identical.
- Different types of noise need different transformations

# Motivation: noise reduction



- Even multiple images of the same static scene will not be identical.
- How could we reduce the noise, i.e., give an estimate of the true intensities?
- **What if there's only one image?**

# Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise

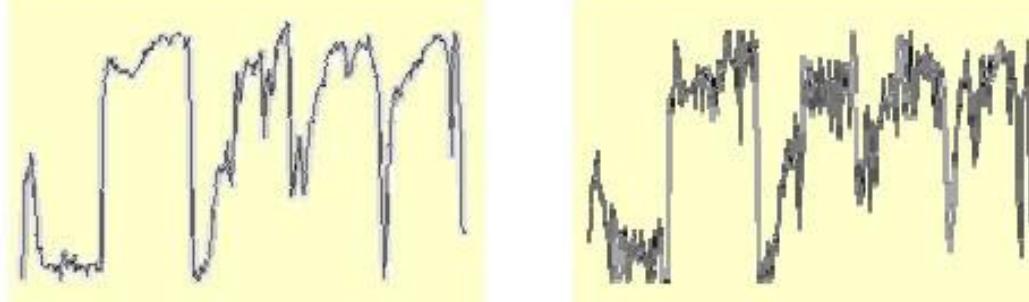
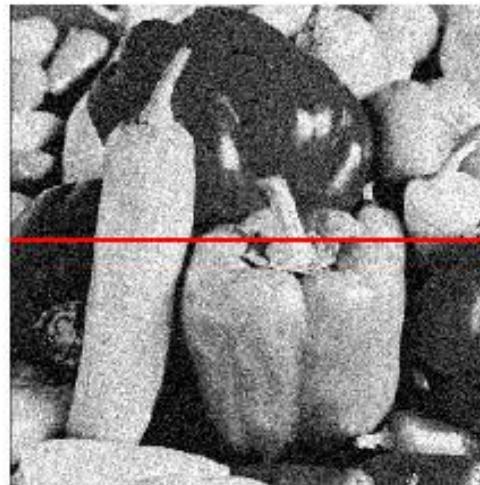


Impulse noise



Gaussian noise

# Gaussian noise



$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:  
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

```
>> randn(noise, 0, sigma);  
>> output = im + noise;
```

What is impact of the sigma?

Fig: M. Hebert

# First attempt at a solution

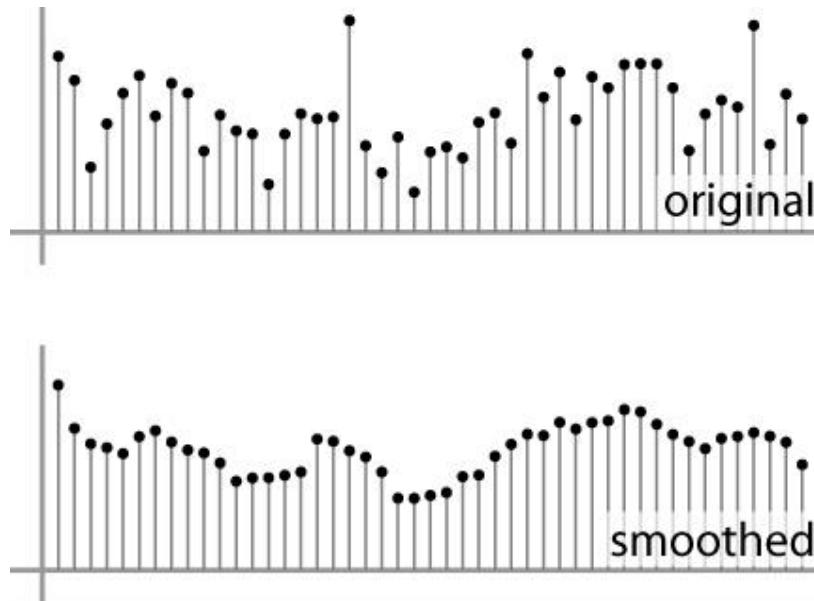
---

- Let's replace each pixel with an average of all the values in its neighborhood
- Assumptions: ([Why?](#))
  - **Expect pixels to be like their neighbors**
  - **Expect noise processes to be independent from pixel to pixel**

# First attempt at a solution

---

- Let's replace each pixel with an average of all the values in its neighborhood
- Moving average in 1D:

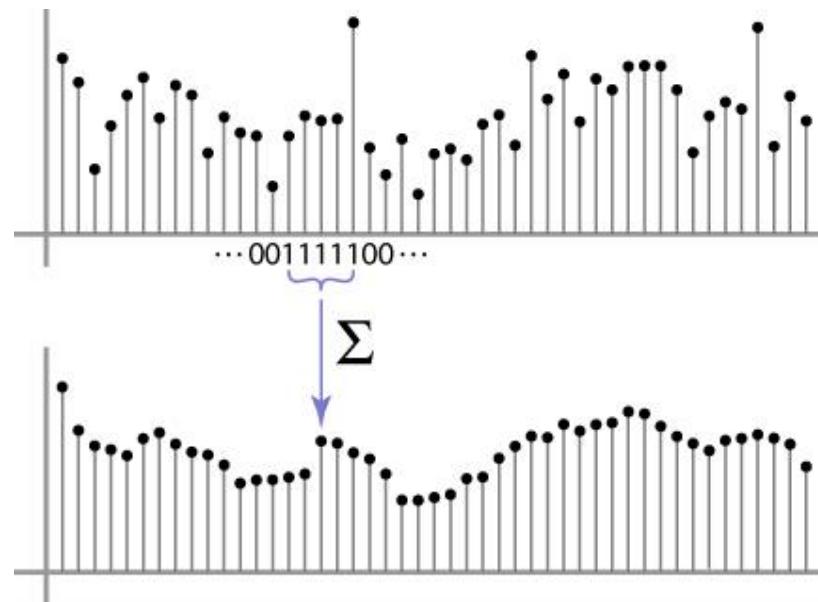


# Weighted Moving Average

---

Can add weights to our moving average

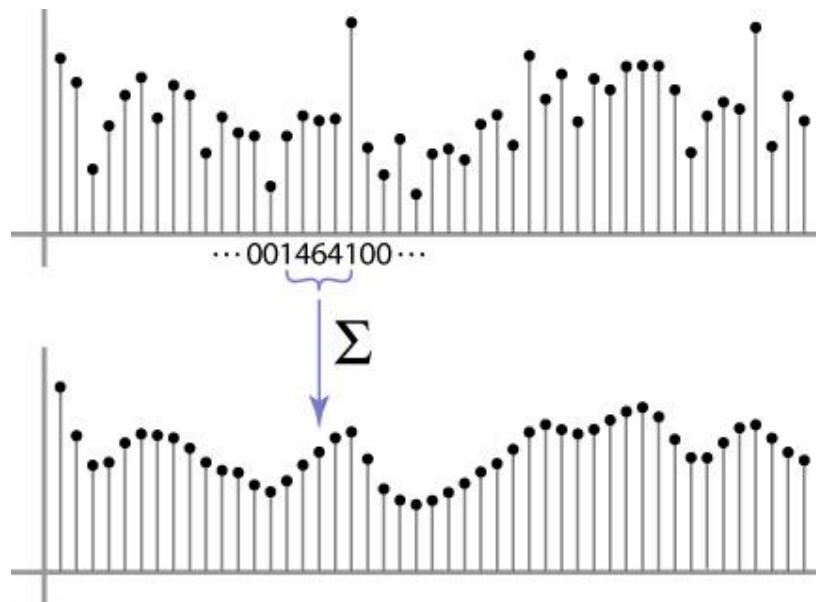
*Weights* [1, 1, 1, 1, 1] / 5



# Weighted Moving Average

---

Non-uniform weights  $[1, 4, 6, 4, 1] / 16$



# Moving Average In 2D

---

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

			0							

# Moving Average In 2D

---

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

0	10									

# Moving Average In 2D

---

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

			0	10	20					

# Moving Average In 2D

---

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$


0    10    20    30

# Moving Average In 2D

---

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$


# Moving Average In 2D

---

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

# Correlation mask

Say the averaging window size is  $2k+1 \times 2k+1$ :

$$G[i, j] = \underbrace{\frac{1}{(2k+1)^2}}_{\text{Attribute uniform weight to each pixel}} \underbrace{\sum_{u=-k}^k \sum_{v=-k}^k F[i+u, j+v]}_{\text{Loop over all pixels in neighborhood around image pixel } F[i,j]}$$

Now generalize to allow **different weights** depending on neighboring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{Non-uniform weights}} F[i+u, j+v]$$

# Correlation mask

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

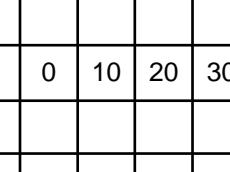
This is called **cross-correlation**, denoted  $G = H \otimes F$

Filtering an image: replace each pixel with a linear combination of its neighbors.

The filter “**kernel**” or “**mask**”  $H[u, v]$  is the prescription for the weights in the linear combination.

# Averaging mask

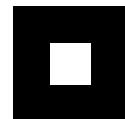
- What values belong in the kernel  $H$  for the moving average example?



A 10x10 grid with horizontal and vertical axes. The horizontal axis is labeled with values 0, 10, 20, and 30. The vertical axis has 10 unlabeled grid lines. A red square box highlights the cell at the intersection of the 4th horizontal line and the 4th vertical line from the top-left corner, which contains the value 30.

$$G = H \otimes F$$

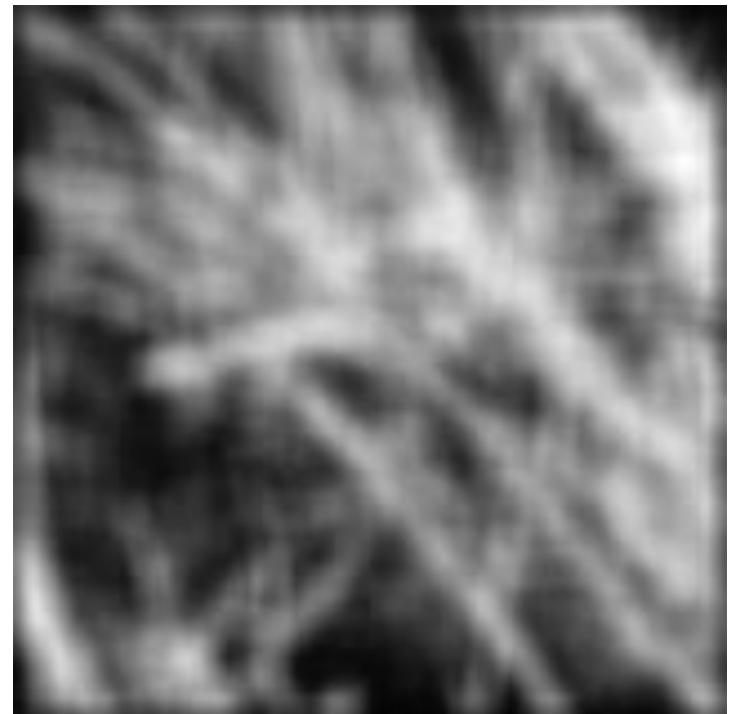
# Smoothing by averaging



depicts box filter:  
white = high value, black = low value



original



filtered

What if the filter size was  $5 \times 5$  instead of  $3 \times 3$ ?

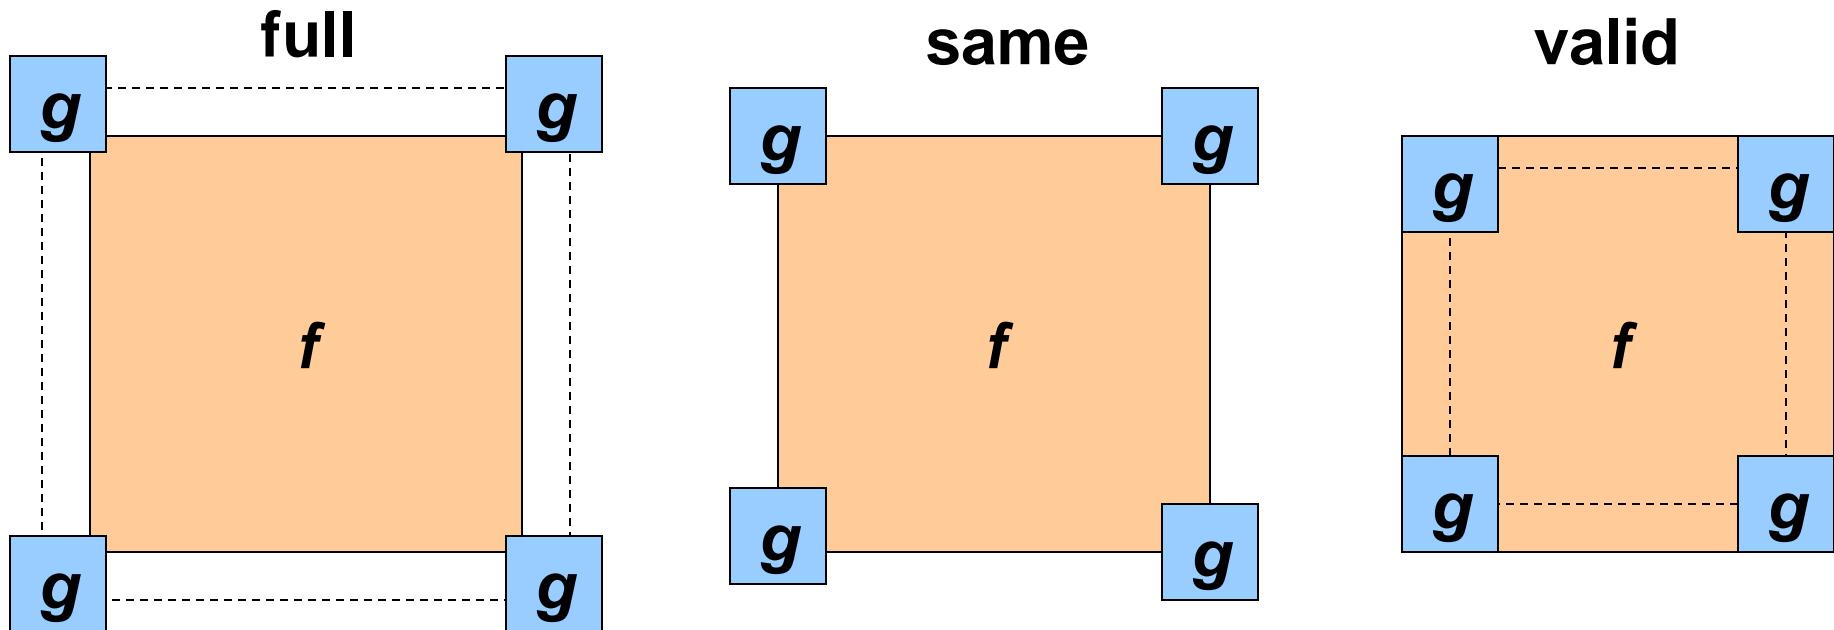
# Boundary issues

---

What is the size of the output?

- Output size / “shape” options

- *shape* = ‘full’: output size is sum of sizes of f and g
- *shape* = ‘same’: output size is same as f
- *shape* = ‘valid’: output size is difference of sizes of f and g

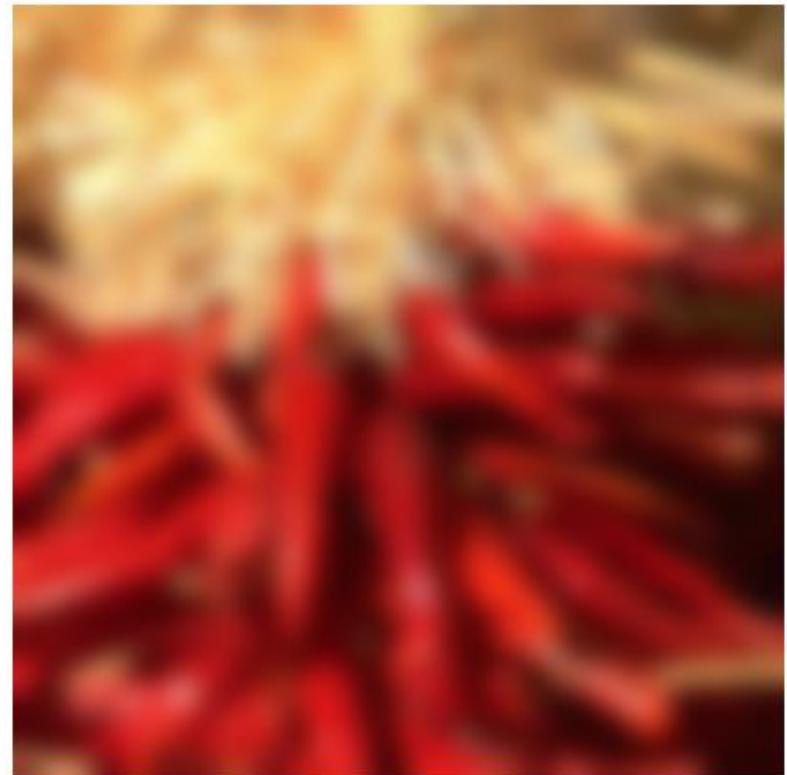


# Boundary issues

---

What about near the edge?

- the filter window falls off the edge of the image
- need to extrapolate
- methods:
  - clip filter (black)
  - wrap around
  - copy edge
  - reflect across edge



# Boundary issues

---

## What about near the edge?

- the filter window falls off the edge of the image
- need to extrapolate
- Types (OpenCV):
  - \* *BORDER\_REPLICATE*: aaaaaa | abcdefgh | hhhhhh
  - \* *BORDER\_REFLECT*: fedcba | abcdefgh | hgfedcb
  - \* *BORDER\_REFLECT\_101*: gfedcb | abcdefgh| gfedcba
  - \* *BORDER\_WRAP*: cdefgh | abcdefgh | abcdefg
  - \* *BORDER\_CONSTANT*: iiиии | abcdefgh | iiиии

# Gaussian mask

- What if we want nearest neighboring pixels to have the most influence on the output?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

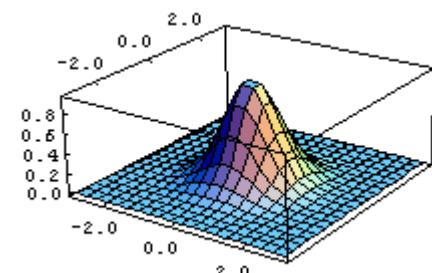
$F[x, y]$

$$\frac{1}{16} \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix}$$

$H[u, v]$

This kernel is an approximation of a 2d Gaussian function:

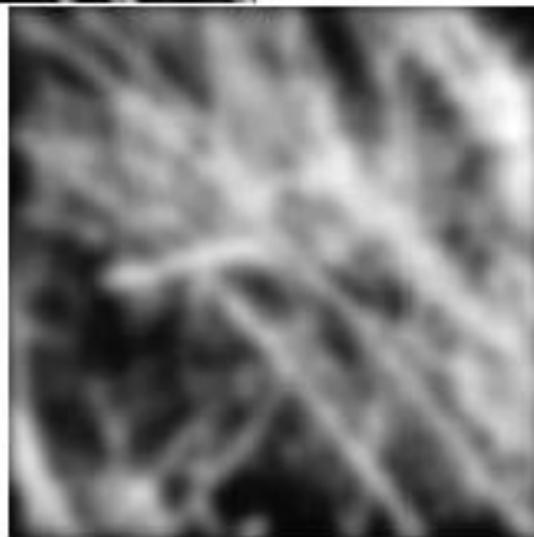
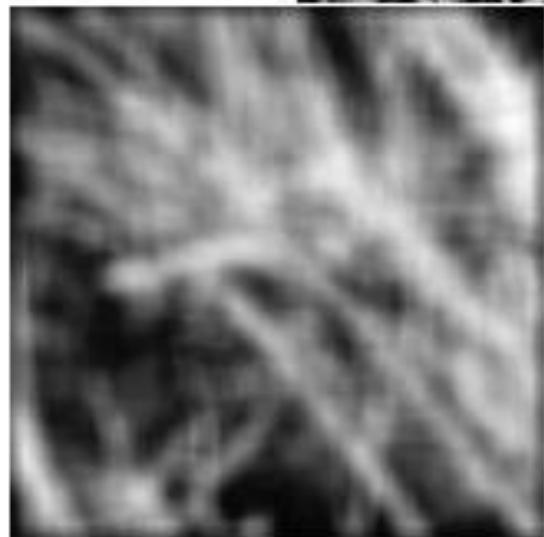
$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



- Removes high-frequency components from the image (“low-pass filter”).

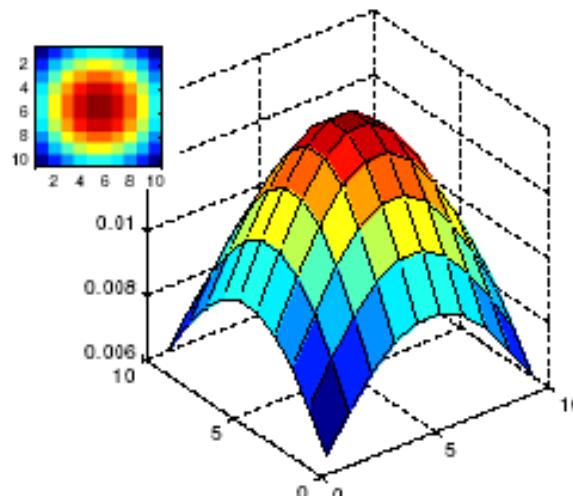
# Gaussian vs. box mask

---

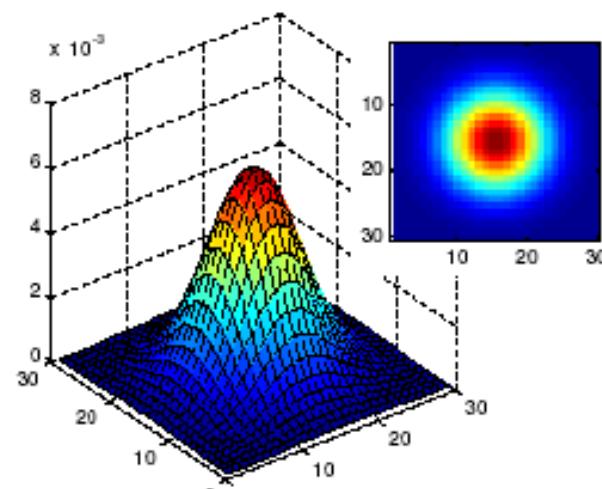


# Gaussian masks

- What parameters matter here?
- **Size** of kernel or mask
  - Note, Gaussian function has infinite support, but discrete filters use finite kernels



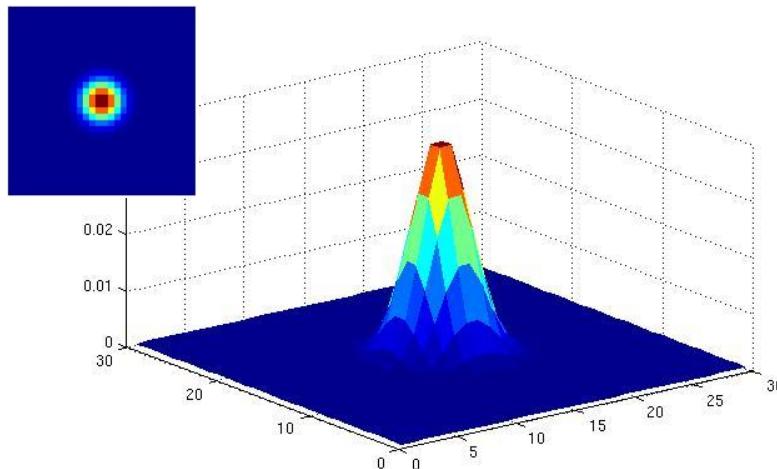
$\sigma = 5$  with  
10 x 10  
kernel



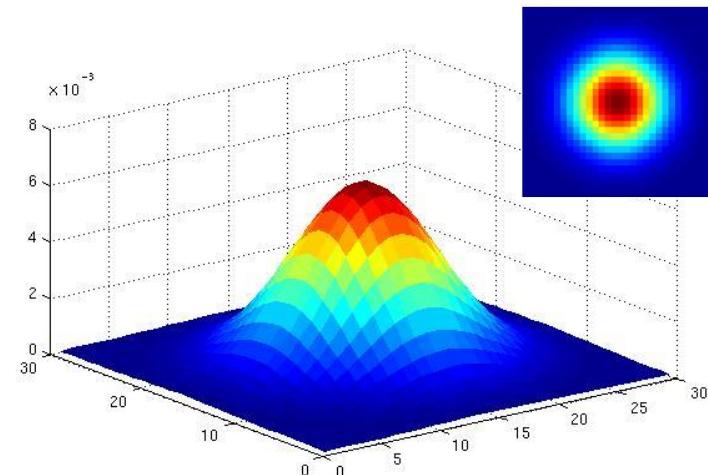
$\sigma = 5$  with  
30 x 30  
kernel

# Gaussian masks

- What parameters matter here?
- **Variance** of Gaussian: determines extent of smoothing



$\sigma = 2$  with  
30 x 30  
kernel



$\sigma = 5$  with  
30 x 30  
kernel

# OpenCV

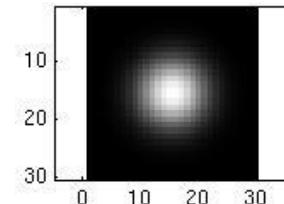
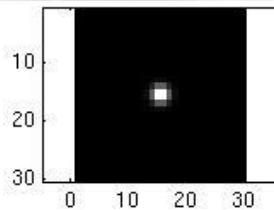
```
>> hsize = 10;  
>> sigma = 5;  
>> cv2.getGaussianKernel(hsize, sigma)  
>> cv2.imshow(kernel);   
  
>> cv2.filter2D(im,outim,kernel); #correlation  
>> cv2.imshow(outim);
```



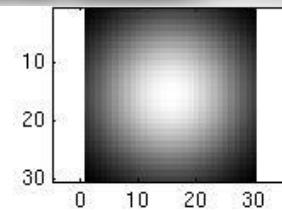
**outim**

# Smoothing with a Gaussian

Parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



• • •



```
for sigma in np.arange(1, 10, 3)
    cv2.getGaussianKernel(hsize, sigma)
    cv2.filter2D(im,outim,kernel); #correlation
    cv2.imshow(outim);
    cv2.waitKey()
```

# Properties of smoothing masks

- Smoothing
  - Values positive
  - Sum to 1 → constant regions same as input
  - Amount of smoothing proportional to mask size
  - **Remove “high-frequency” components; “low-pass” filter**

# Convolution

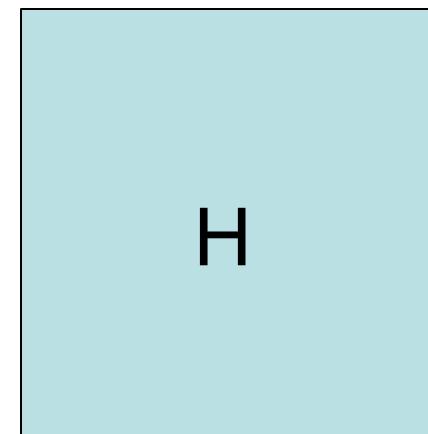
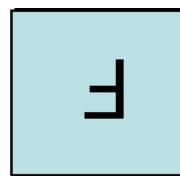
- Convolution:
  - Flip the filter in both dimensions (bottom to top, right to left)
  - Then apply cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$



*Notation for  
convolution  
operator*



# Convolution vs. correlation

## Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i - u, j - v]$$

$$G = H \star F$$

## Cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i + u, j + v]$$

$$G = H \otimes F$$

For a Gaussian or box filter, how will the outputs differ?

If the input is an impulse signal, how will the outputs differ?

# Predict the outputs using correlation filtering



$$\text{Input} * \begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix} = ?$$



$$\text{Input} * \begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{matrix} = ?$$

$$\text{Input} * \begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix} - \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix} = ?$$

# Practice with linear filters

---



Original

0	0	0
0	1	0
0	0	0

?

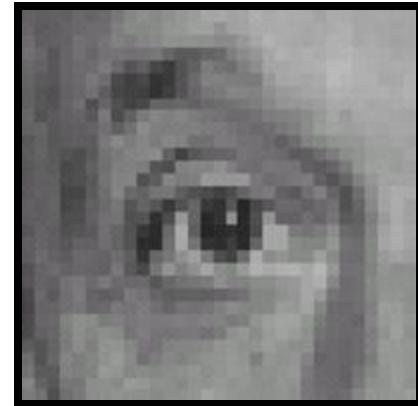
# Practice with linear filters

---



**Original**

0	0	0
0	1	0
0	0	0



**Filtered  
(no change)**

# Practice with linear filters

---



Original

0	0	0
0	0	1
0	0	0

?

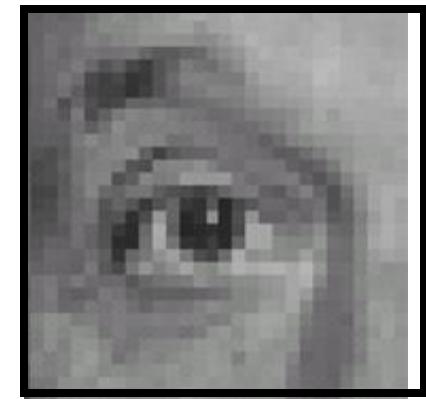
# Practice with linear filters

---



**Original**

0	0	0
0	0	1
0	0	0



**Shifted left  
by 1 pixel  
with  
correlation**

# Practice with linear filters

---



$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

?

Original

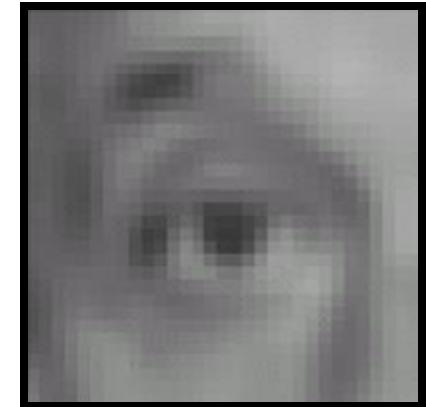
# Practice with linear filters

---



Original

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



Blur (with a  
box filter)

# Practice with linear filters

---



$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

-

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

?

Original

# Practice with linear filters

---

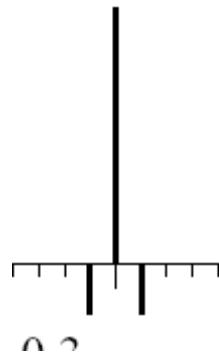


Original

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

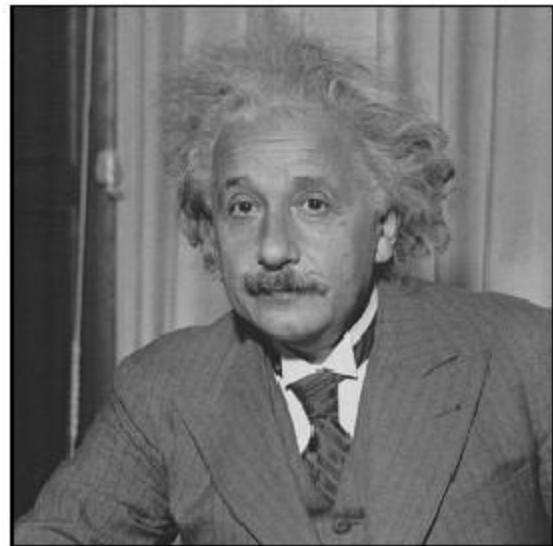
-

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

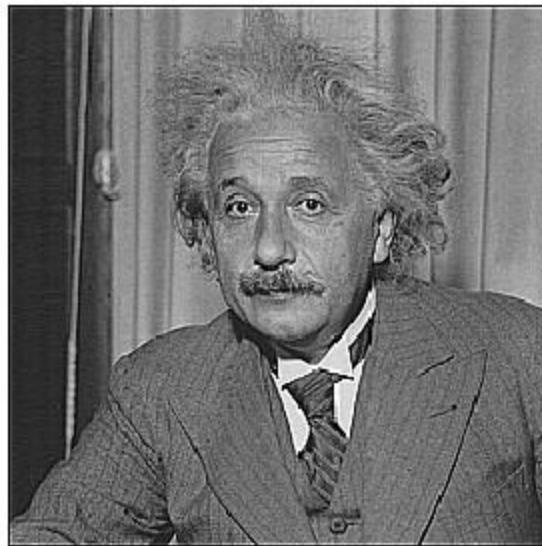


Sharpening filter:  
accentuates differences  
with local average

# Filtering examples: sharpening



before



after

# Sharpening

---

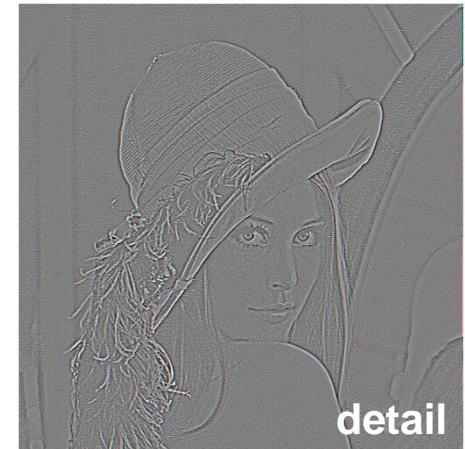
What does blurring take away?



-



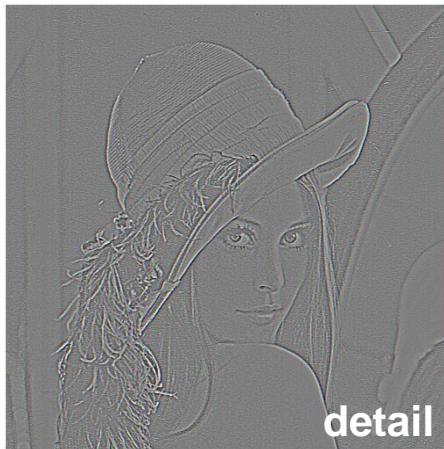
=



Let's add it back:



+



=



# Properties of convolution

- **Shift invariant:**
  - Operator behaves the same everywhere, i.e. the value of the output depends on the pattern in the image neighborhood, not the position of the neighborhood.
- **Superposition:**
  - $h * (f_1 + f_2) = (h * f_1) + (h * f_2)$

# Properties of convolution

- Commutative:

$$f * g = g * f$$

- Associative

$$(f * g) * h = f * (g * h)$$

- Distributes over addition

$$f * (g + h) = (f * g) + (f * h)$$

- Scalars factor out

$$kf * g = f * kg = k(f * g)$$

- Identity:

$$\text{unit impulse } e = [\dots, 0, 0, 1, 0, 0, \dots]. \quad f * e = f$$

# Separability

- In some cases, filter is separable, and we can factor into two steps:
  - Convolve all rows
  - Convolve all columns

# Separability

- In some cases, filter is separable, and we can factor into two steps: e.g.,

$$\begin{array}{c} \mathbf{g} \\ \mathbf{f} \\ \mathbf{h} \end{array} \quad \begin{array}{c} \begin{matrix} 1 & 2 & 1 \end{matrix} \\ \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} \\ \begin{matrix} 11 & & \\ & 18 & \\ & 18 & \end{matrix} \end{array}$$

What is the computational complexity advantage for a separable filter of size  $k \times k$ , in terms of number of operations per output pixel?

$$\mathbf{f} * (\mathbf{g} * \mathbf{h}) = (\mathbf{f} * \mathbf{g}) * \mathbf{h}$$

# Why is separability useful?

---

- Separability means that a 2D convolution can be reduced to two 1D convolutions (one among rows and one among columns)
- What is the complexity of filtering an  $n \times n$  image with an  $m \times m$  kernel?
  - $O(n^2 m^2)$
- What if the kernel is separable?
  - $O(n^2 m)$

# Gaussian filters

---

- Remove high-frequency components from the image (*low-pass filter*)
- Convolution with self is another Gaussian
  - So can smooth with small- $\sigma$  kernel, repeat, and get same result as larger- $\sigma$  kernel would have
  - Convolving two times with Gaussian kernel with std. dev.  $\sigma$  is same as convolving once with kernel with std. dev.  $\sigma\sqrt{2}$
- **Separable kernel**
  - Factors into product of two 1D Gaussians
  - Discrete example:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

# Separability of the Gaussian filter

---

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

# Effect of smoothing filters

5x5

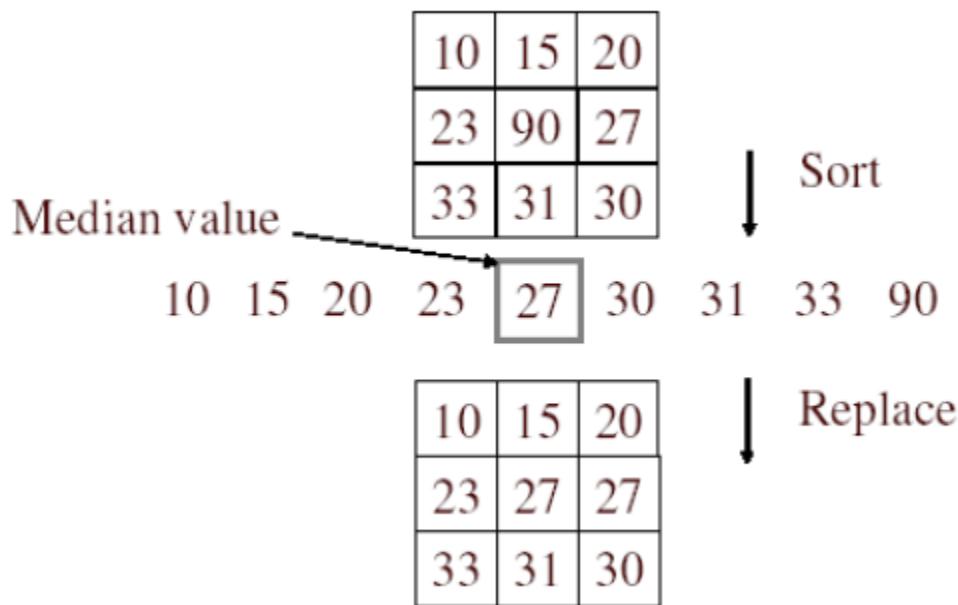


Additive Gaussian noise



Salt and pepper noise

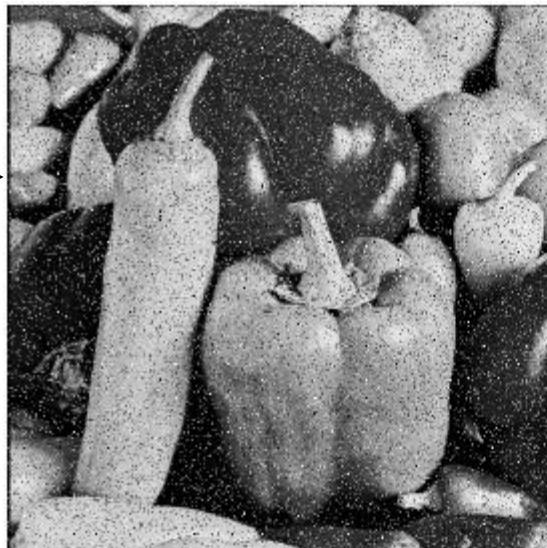
# Median filter



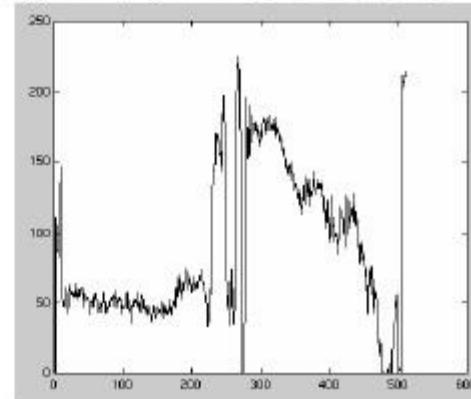
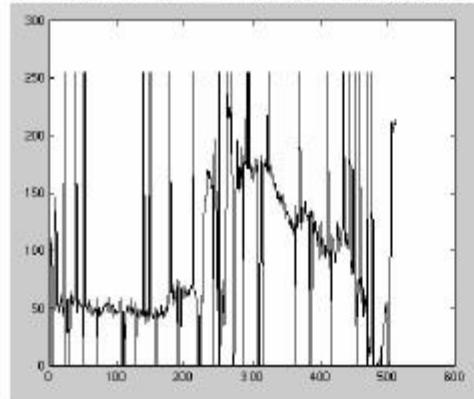
- No new pixel values introduced
- Removes spikes: good for impulse, salt & pepper noise
- Non-linear filter

# Median filter

Salt and pepper noise



Median filtered



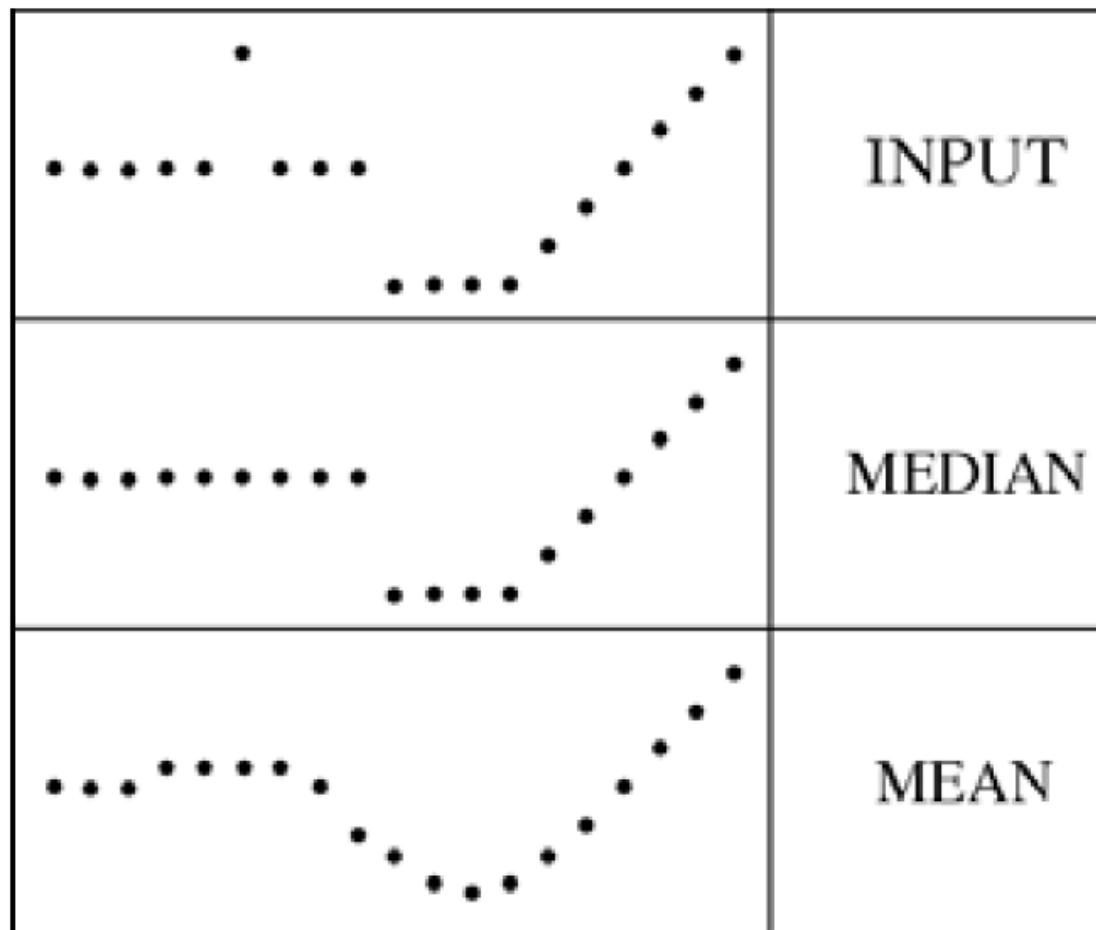
Plots of a row of the image

OpenCV: `cv2.medianBlur(src, ksize[, dst ])->dst`

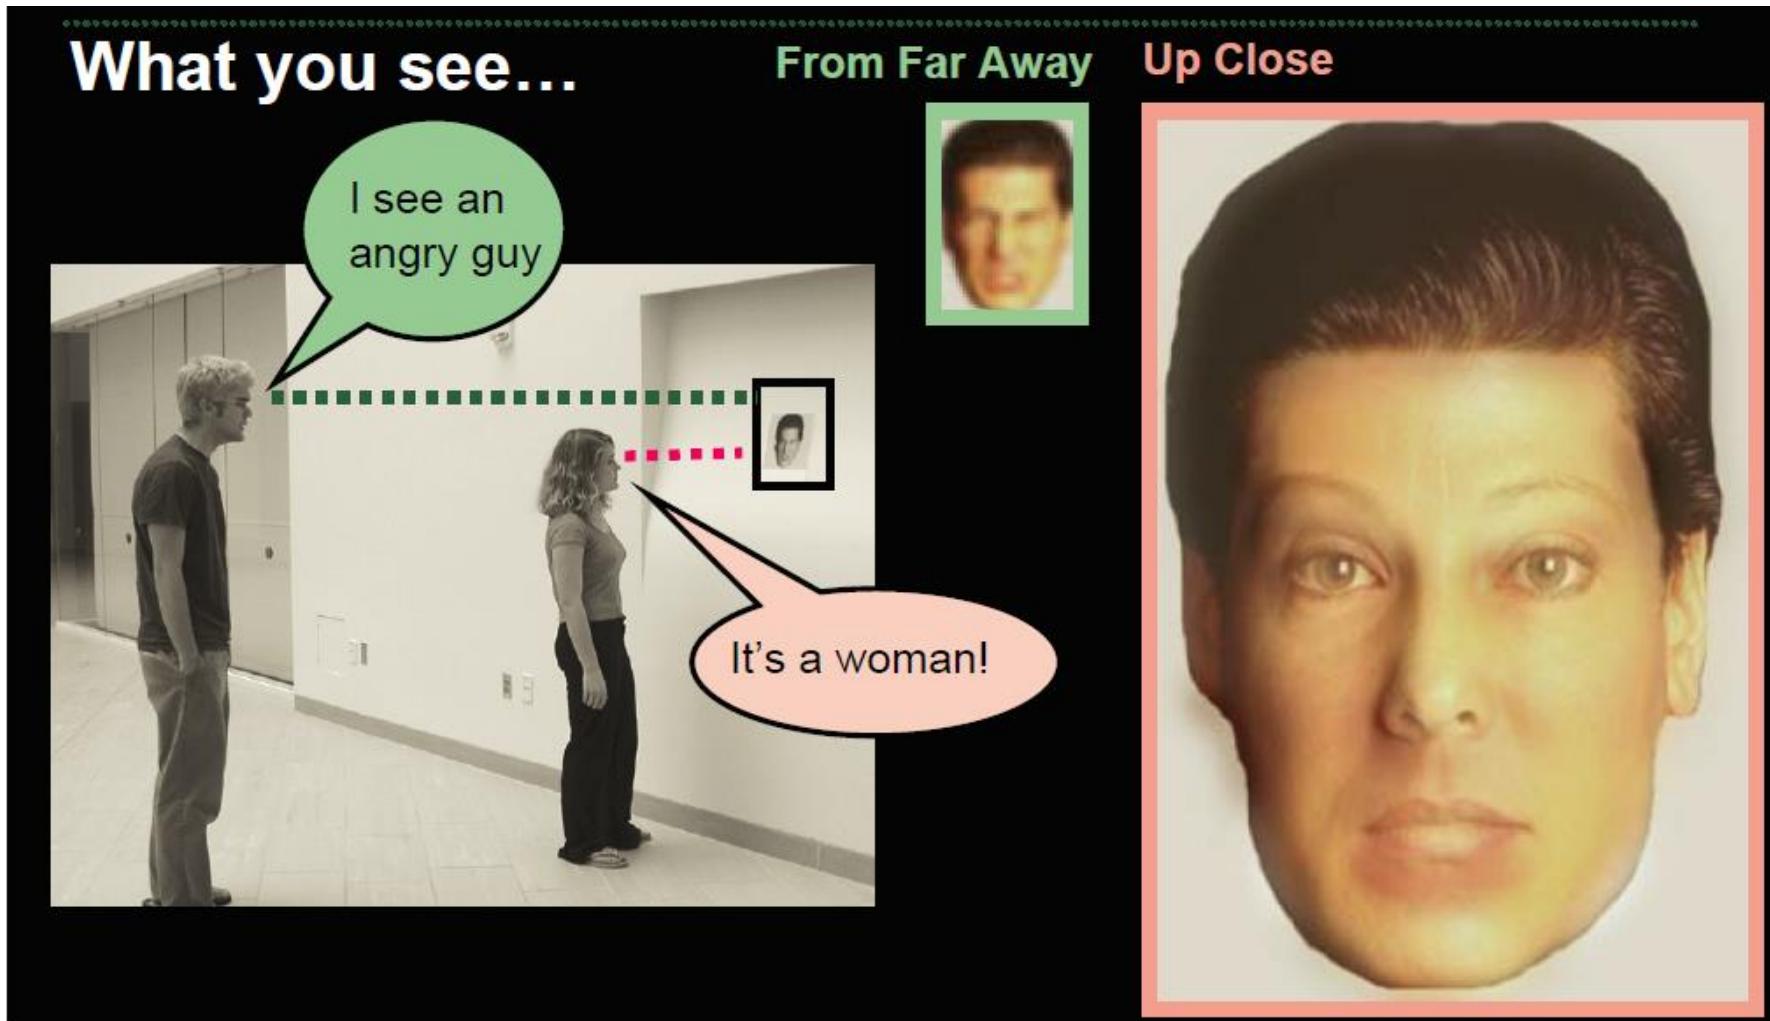
Source: M. Hebert

# Median filter

- Median filter is edge preserving

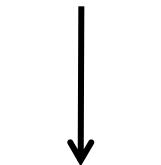


# Filtering application: Hybrid Images

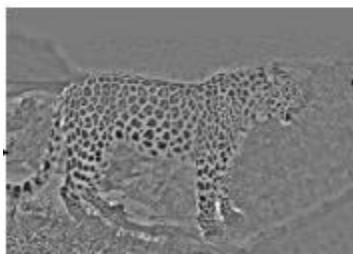
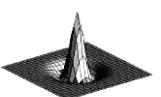
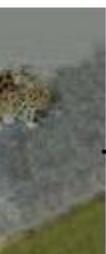


# Application: Hybrid Images

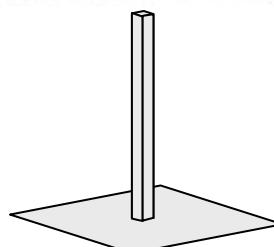
Gaussian Filter



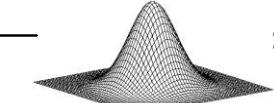
A. Oliva, A. Torralba, P.G. Schyns,  
"Hybrid Images," SIGGRAPH 2006



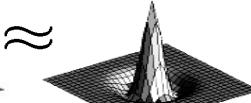
Laplacian Filter



unit impulse

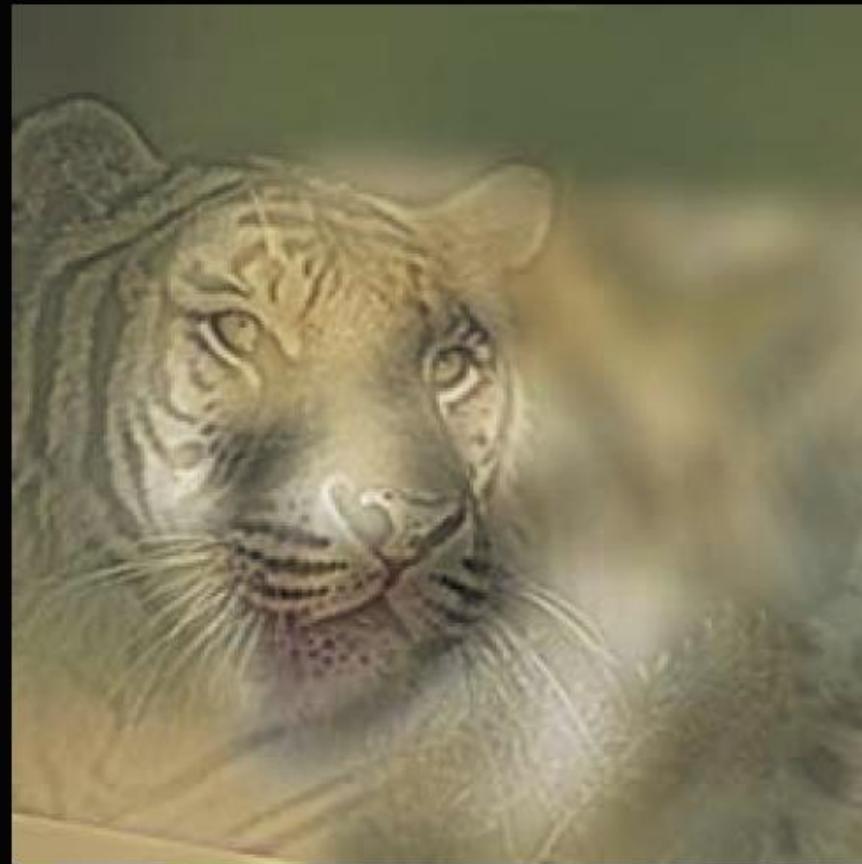


Gaussian



$\approx$  Laplacian of Gaussian





Aude Oliva & Antonio Torralba & Philippe G Schyns, SIGGRAPH 2006

# Changing expression



SIGGRAPH2006

Sad



Surprised



# Summary

- Image “noise”
- Linear filters and convolution useful for
  - Enhancing images (smoothing, removing noise)
    - Box filter
    - Gaussian filter
    - **Impact of scale / width of smoothing filter**
  - Detecting features (next time)
- Separable filters more efficient
- Median filter: a non-linear filter, edge-preserving
- Different frequency bands contribute different information.

# Features: Edge detection

# Edge detection

- **Goal:** map image from 2d array of pixels to a set of curves or line segments or contours.
- **Why?**

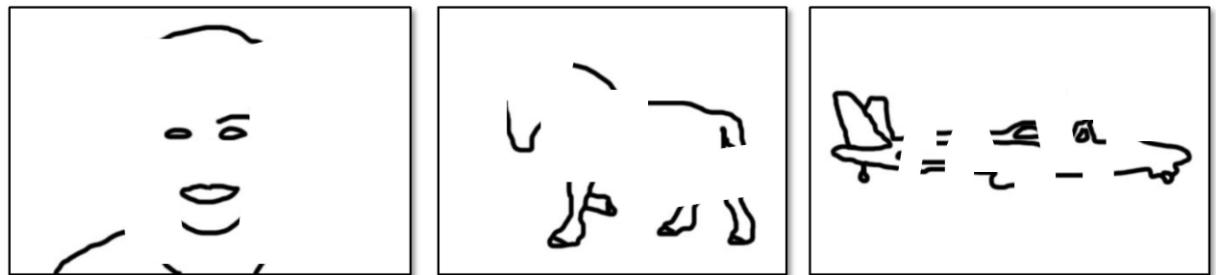
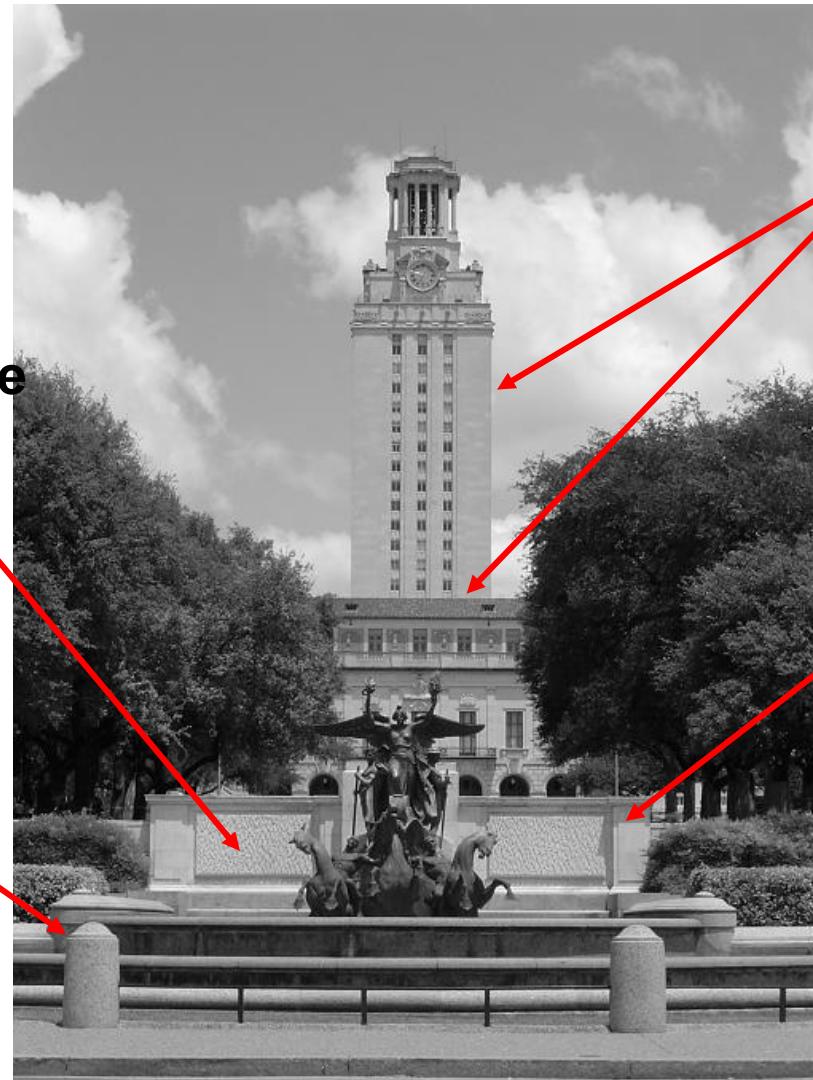


Figure from J. Shotton et al., PAMI 2007

- **Main idea:** look for strong gradients, post-process

# What causes an edge?



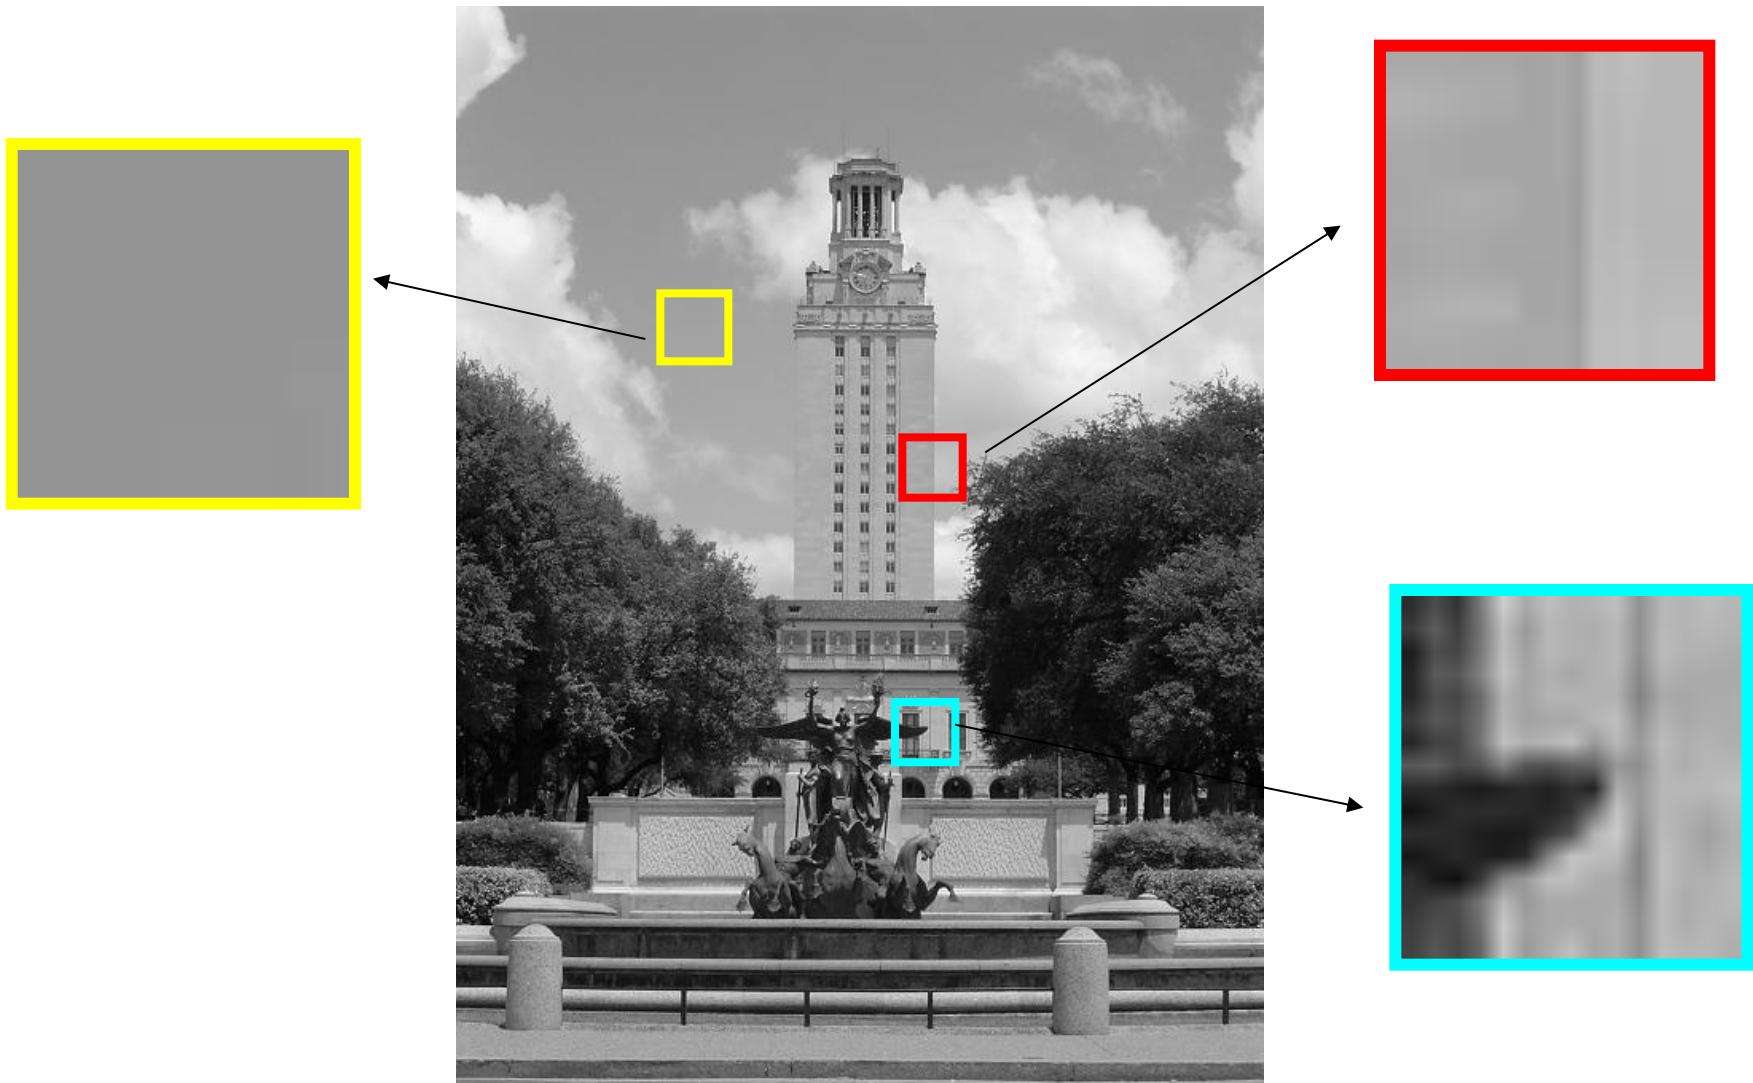
Reflectance  
change:  
appearance  
information, texture

Change in  
surface  
orientation:  
shape

Depth  
discontinuity:  
object boundary

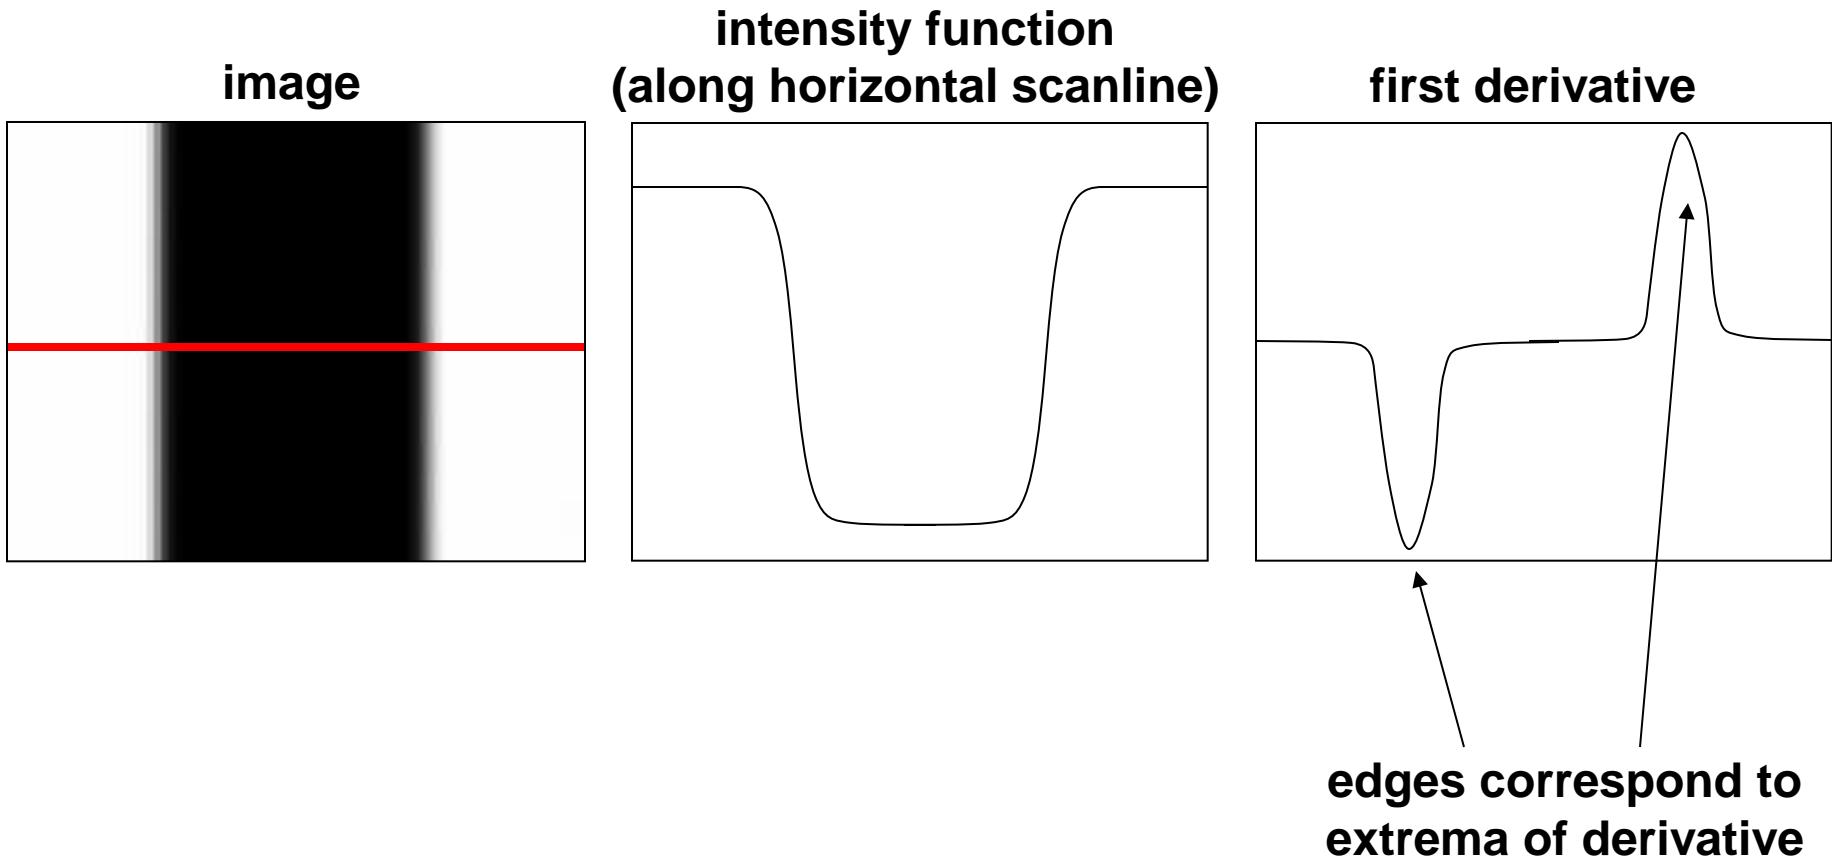
Cast shadows

# Edges/gradients and invariance



# Derivatives and edges

An edge is a place of rapid change in the image intensity function.



# Derivatives with convolution

For 2D function,  $f(x,y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

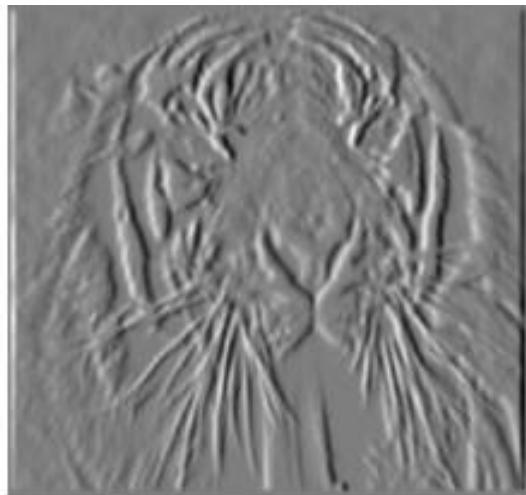
To implement above as convolution, what would be the associated filter?

# Partial derivatives of an image



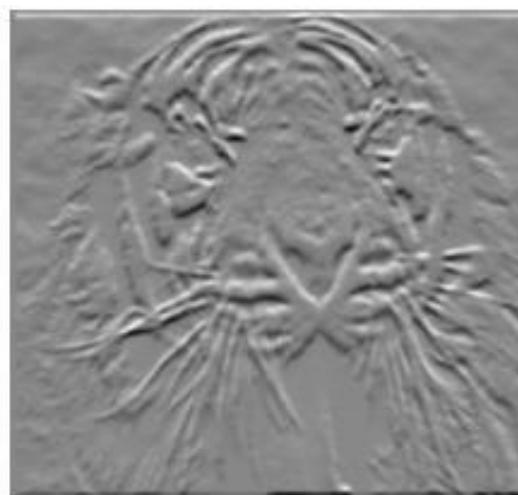
$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

-1	?	1
1	or	-1



Which shows changes with respect to x?

(showing filters for correlation)

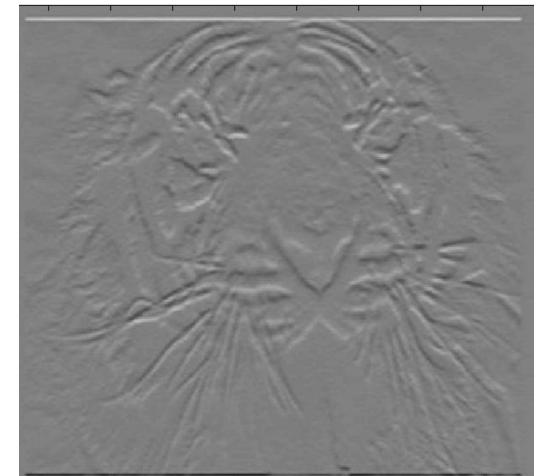
# Assorted finite difference filters

Prewitt:  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

Sobel:  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

Roberts:  $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

```
cv2.Sobel(src, ddepth, dx, dy[, dst[, ksize[, scale[, delta[,  
borderType ]]]]])! ->dst
```



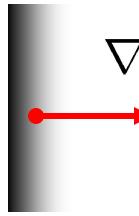
# Image gradient

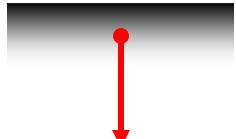
---

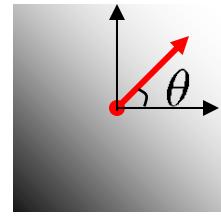
The gradient of an image:

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity


$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The **gradient direction** (orientation of edge normal) is given by:

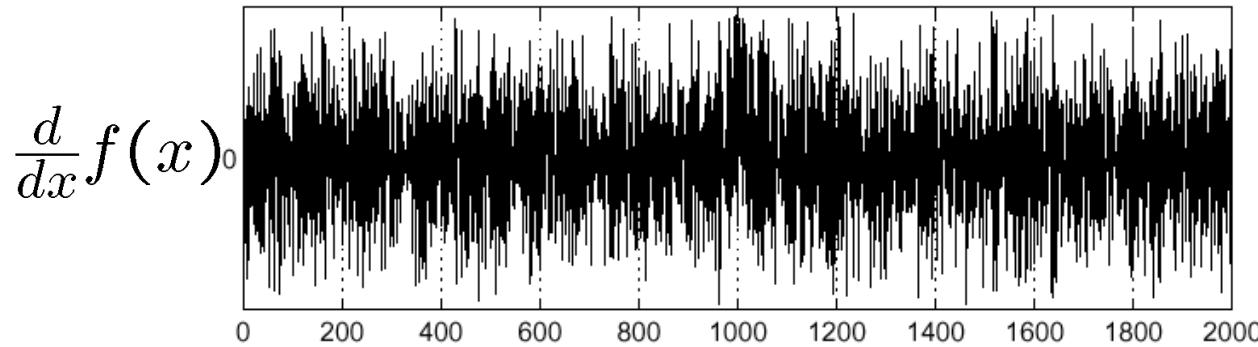
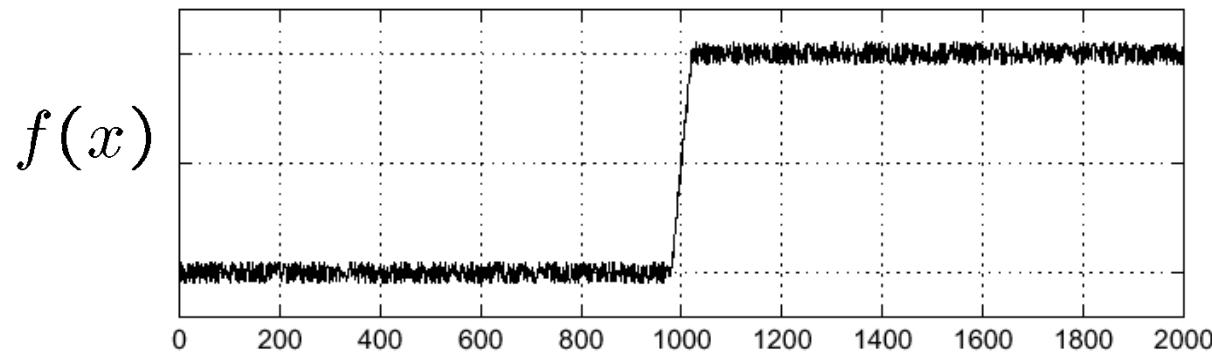
$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

# Effects of noise

---

Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal

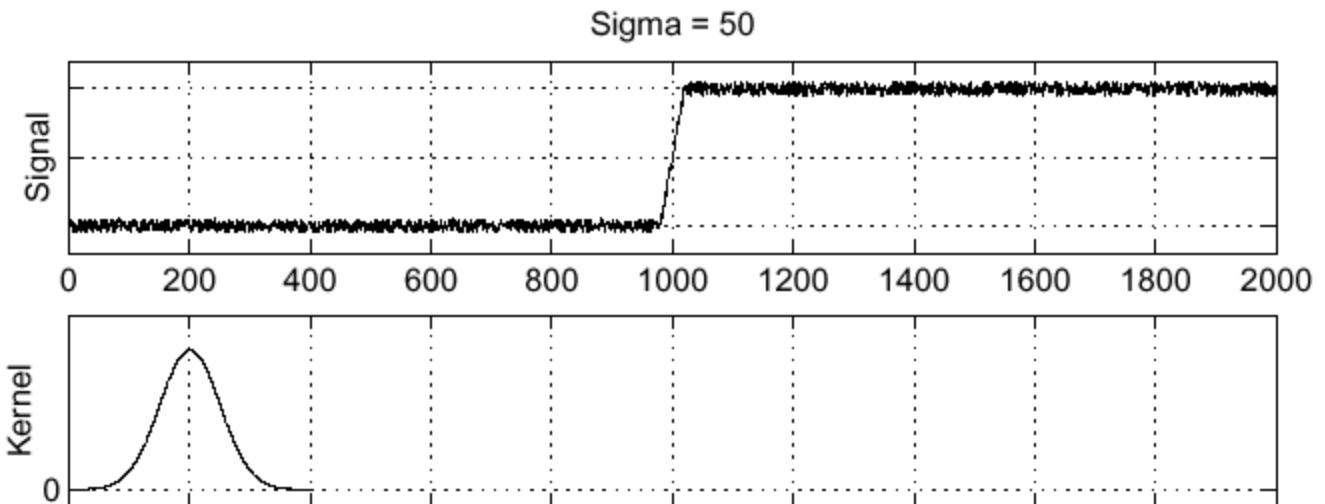


Where is the edge?

# Solution: smooth first

---

$f$



$h$

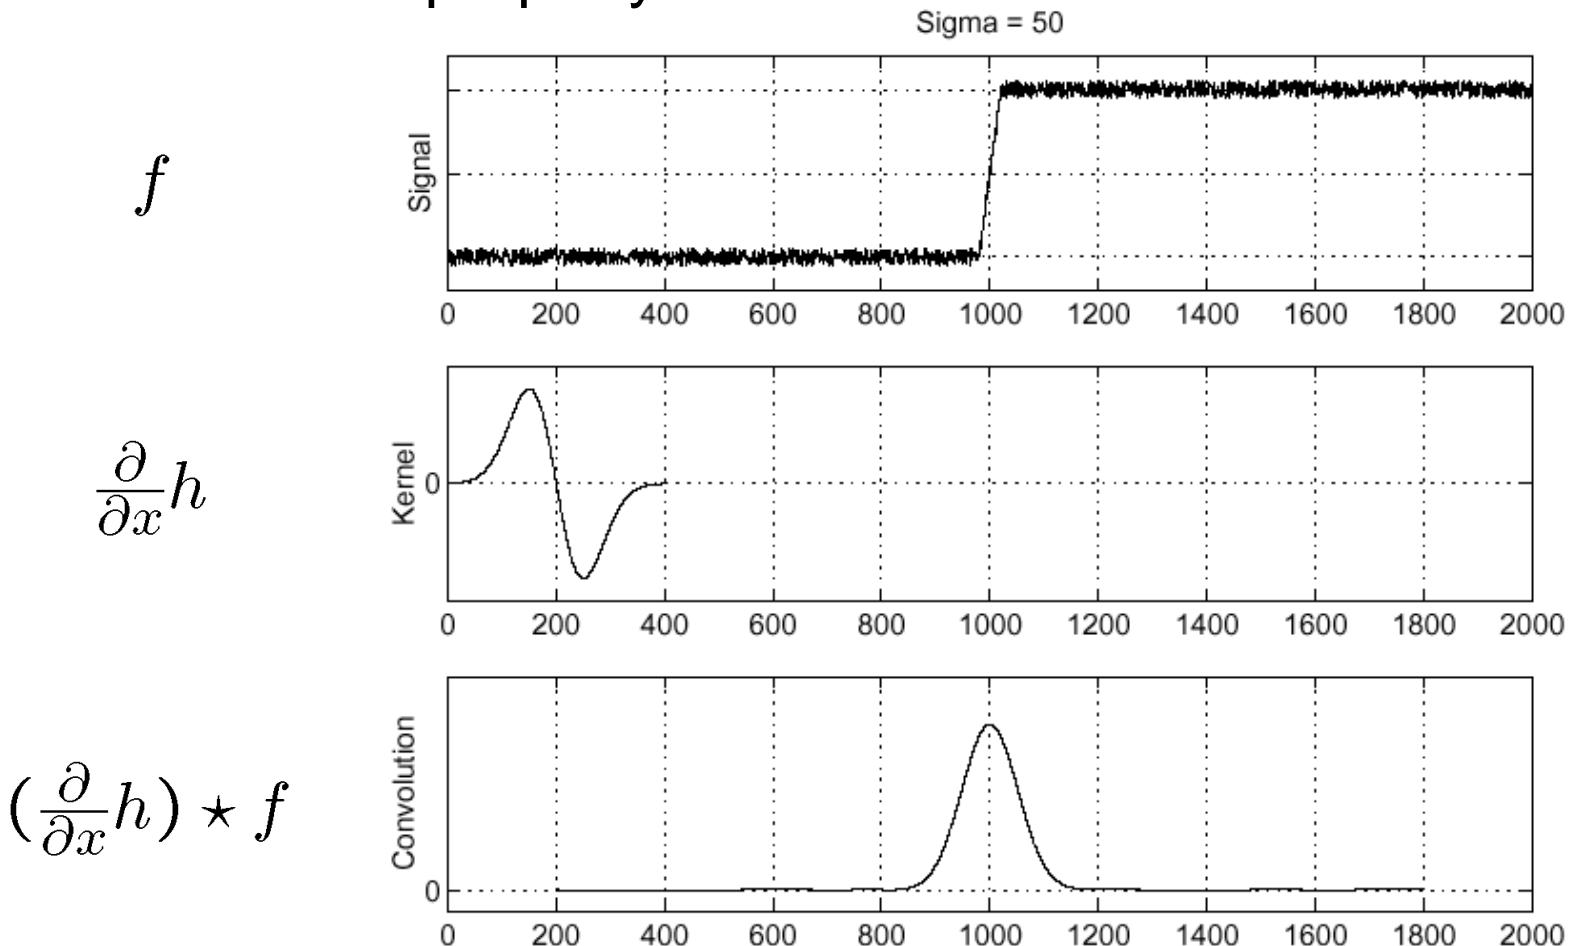
Where is the edge?

Look for peaks in  $\frac{\partial}{\partial x}(h \star f)$

# Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

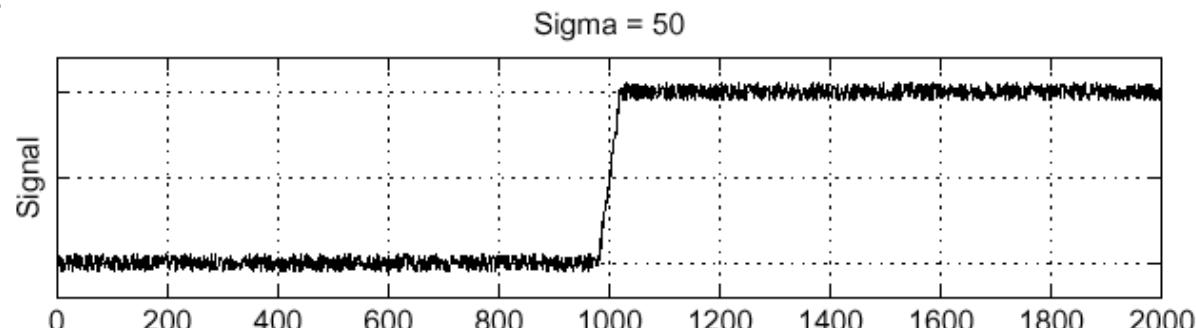
Differentiation property of convolution.



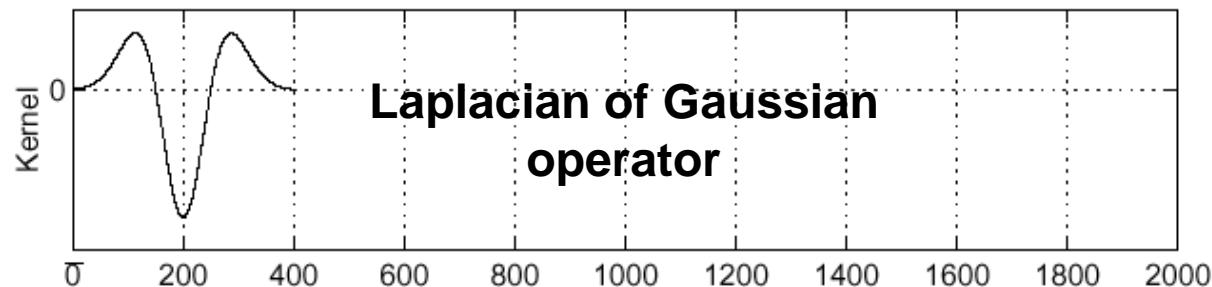
# Laplacian of Gaussian

Consider  $\frac{\partial^2}{\partial x^2}(h \star f)$

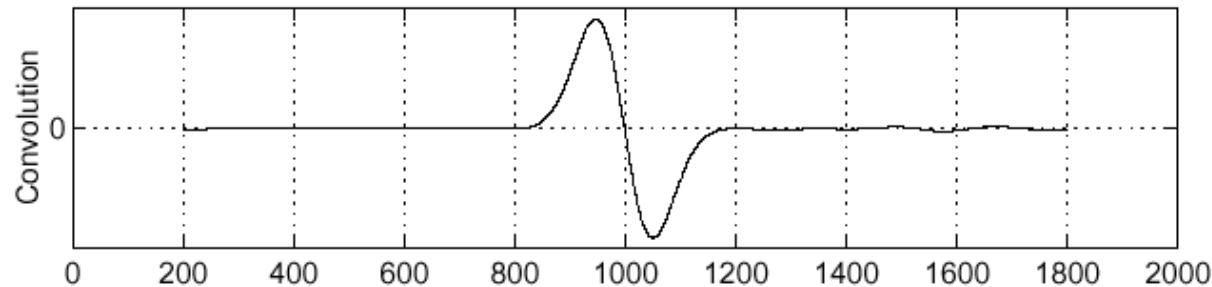
$f$



$\frac{\partial^2}{\partial x^2} h$



$(\frac{\partial^2}{\partial x^2} h) \star f$



Where is the edge?

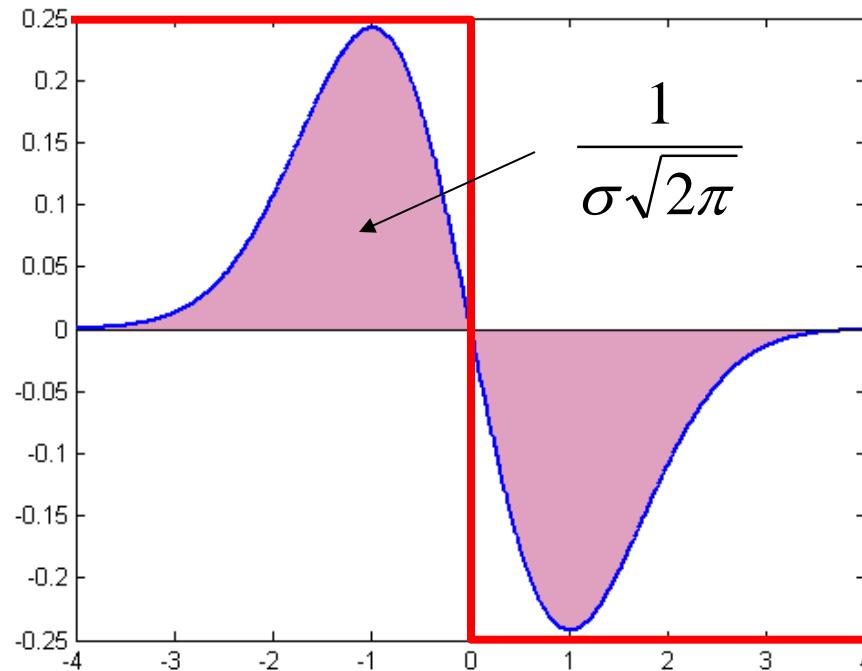
Zero-crossings of bottom graph

Slide credit: Steve Seitz

# Scale normalization

---

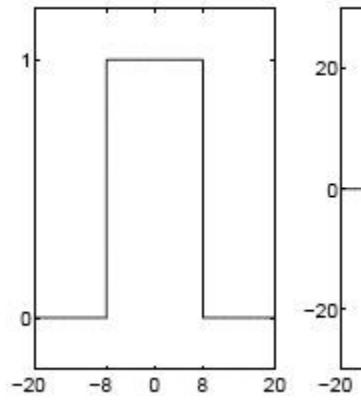
- The response of a derivative of Gaussian filter to a perfect step edge decreases as  $\sigma$  increases



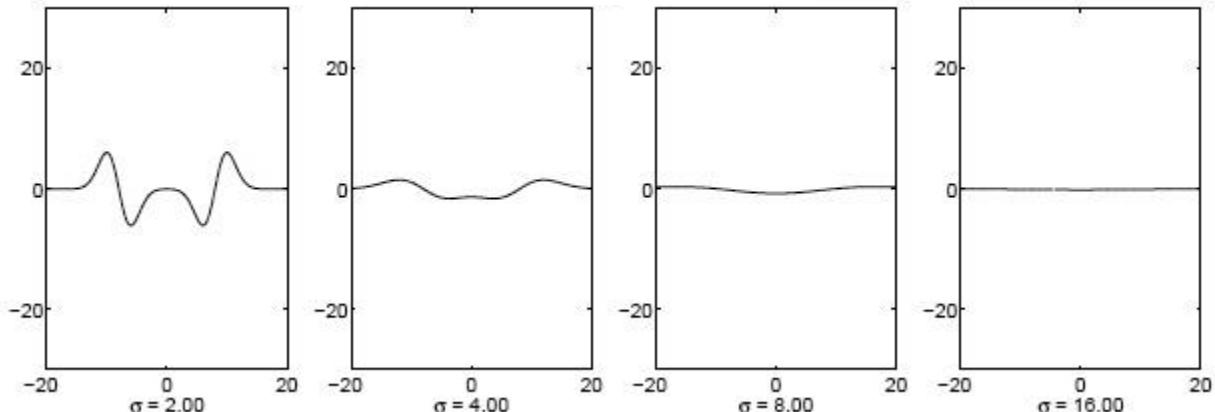
# Effect of scale normalization

---

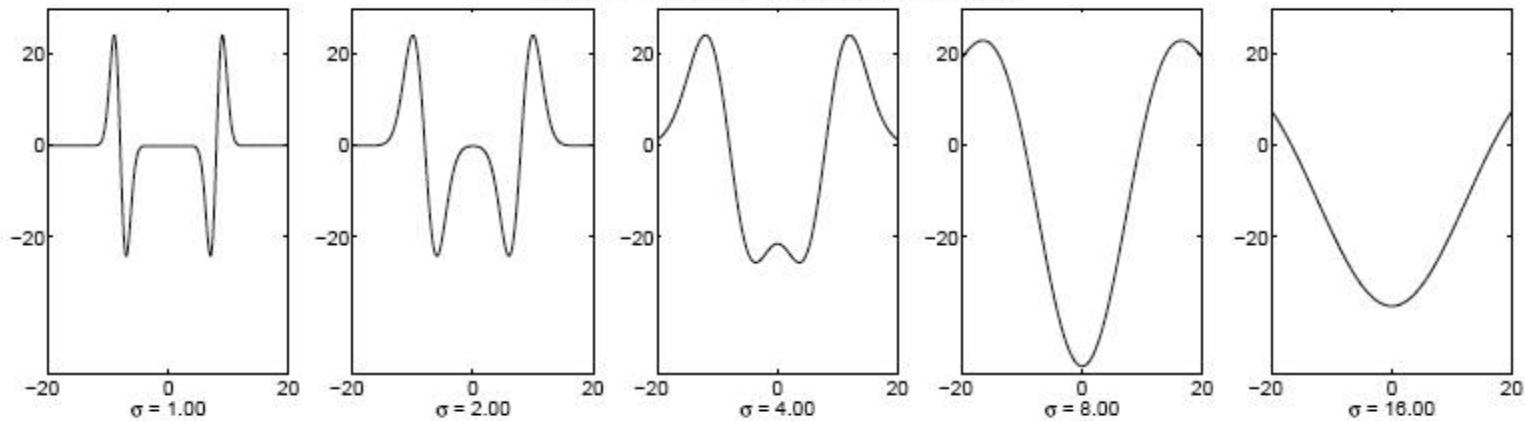
Original signal



Unnormalized Laplacian response



Scale-normalized Laplacian response



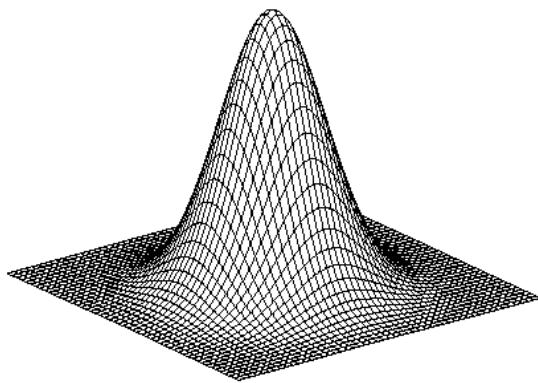
maximum

# Scale normalization

---

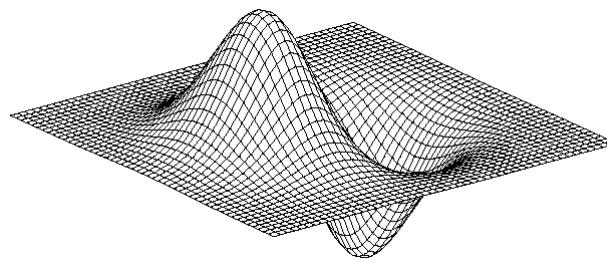
- The response of a derivative of Gaussian filter to a perfect step edge decreases as  $\sigma$  increases
- To keep response the same (scale-invariant), must multiply Gaussian derivative by  $\sigma$
- Laplacian is the second Gaussian derivative, so it must be multiplied by  $\sigma^2$

# 2D edge detection filters



**Gaussian**

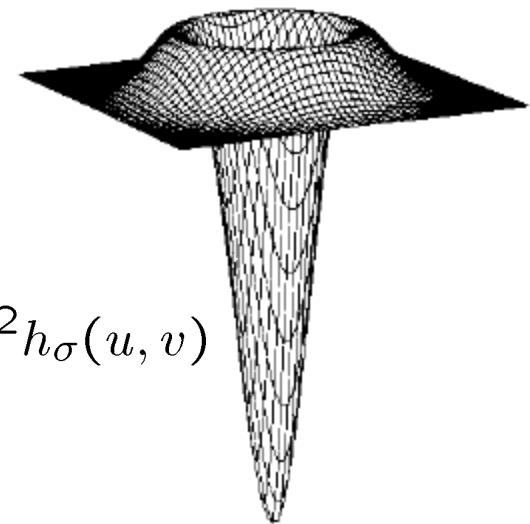
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



**derivative of Gaussian**

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

**Laplacian of Gaussian**



- $\nabla^2$  is the Laplacian operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

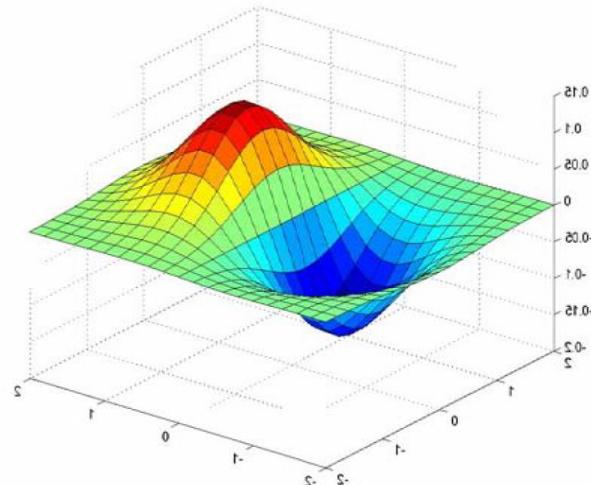
$$\nabla_{\text{norm}}^2 g = \sigma^2 \left( \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

# Derivative of Gaussian filters

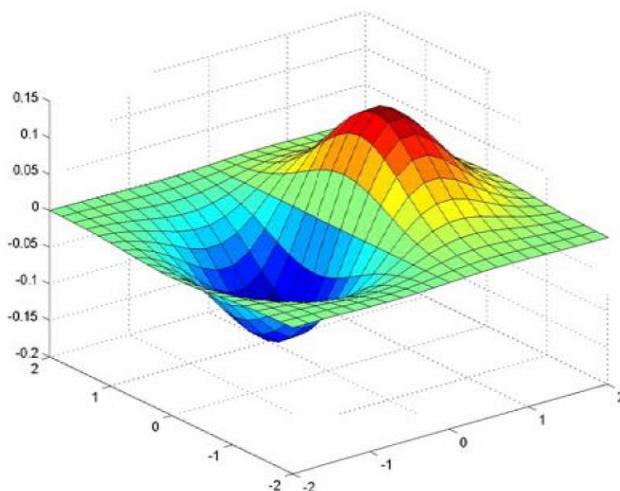
$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$

$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix}$$

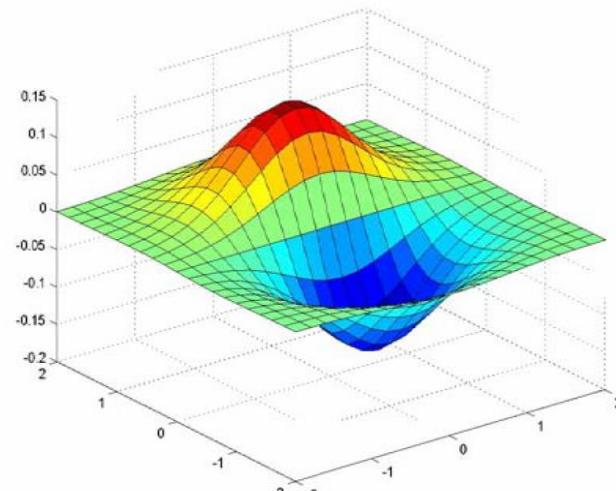
$$\begin{bmatrix} 1 & -1 \end{bmatrix} \otimes \begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix}$$



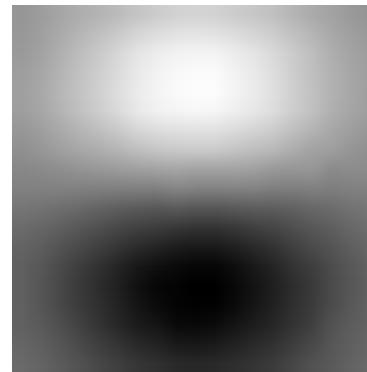
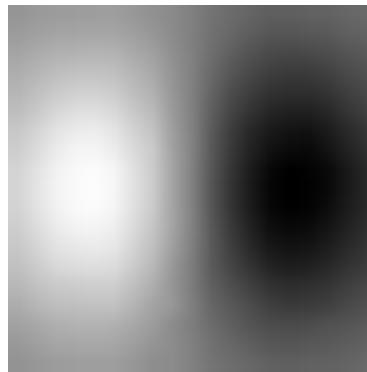
# Derivative of Gaussian filters



x-direction

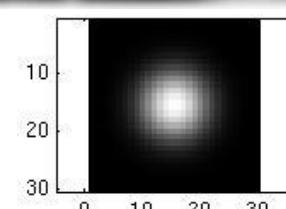
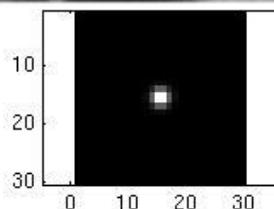


y-direction

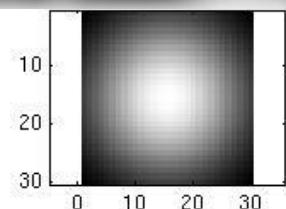


# Smoothing with a Gaussian

Recall: parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



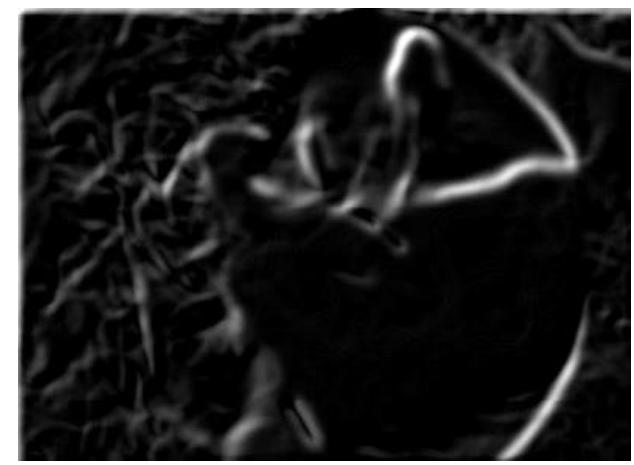
...



# Effect of $\sigma$ on derivatives



$\sigma = 1$  pixel



$\sigma = 3$  pixels

The apparent structures differ depending on Gaussian's scale parameter.

Larger values: larger scale edges detected  
Smaller values: finer features detected

# So, what scale to choose?

It depends what we're looking for.

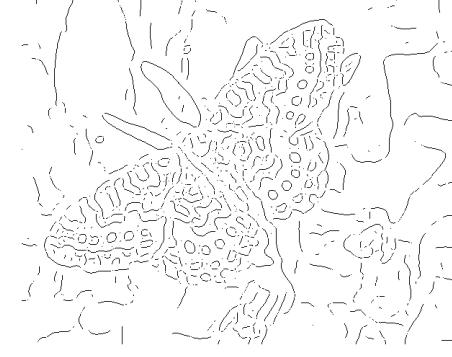


# Mask properties

- Smoothing
  - Values positive
  - Sum to 1 → constant regions same as input
  - Amount of smoothing proportional to mask size
  - Remove “high-frequency” components; “low-pass” filter
- Derivatives
  - Different signs used to get high response in regions of high contrast
  - Sum to zero → no response in constant regions
  - High absolute value at points of high contrast



# Gradients -> edges



Primary edge detection steps:

1. Smoothing: suppress noise
2. Edge enhancement: filter for contrast
3. Edge localization

Determine which local maxima from filter output  
are actually edges vs. noise

- Threshold,
- Thin
- No good results !

# Original image



# Gradient magnitude image



# Thresholding gradient with a lower threshold



# Thresholding gradient with a higher threshold



# Low-level edges vs. perceived contours



Background



Texture



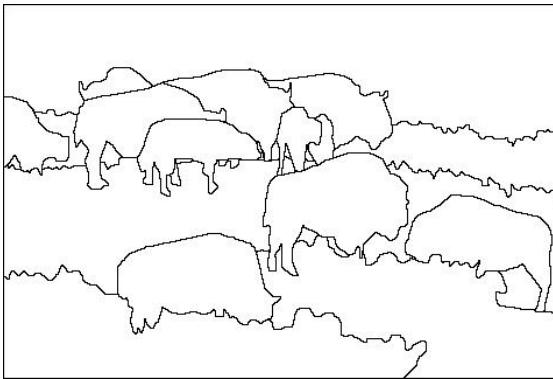
Shadows

# Low-level edges vs. perceived contours

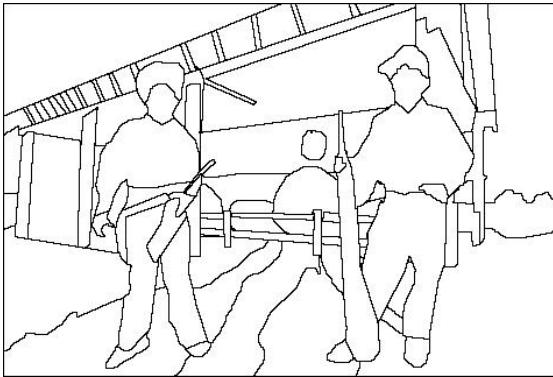
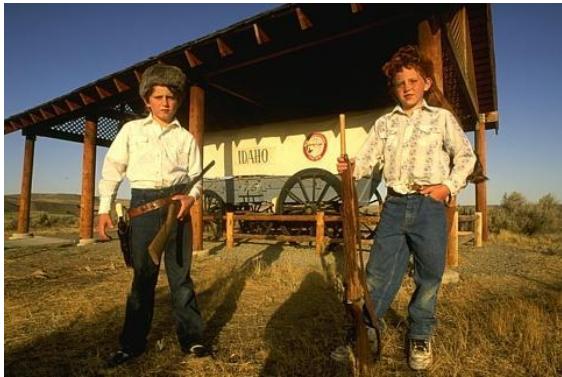
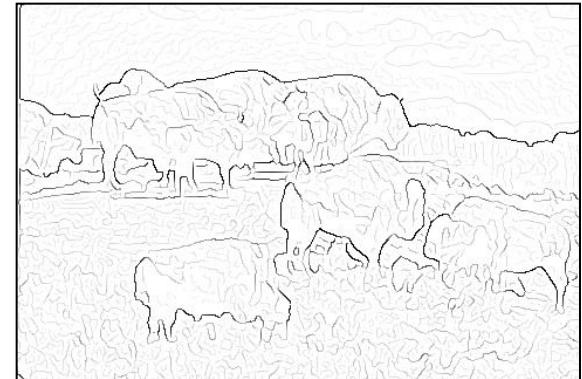
image



human segmentation



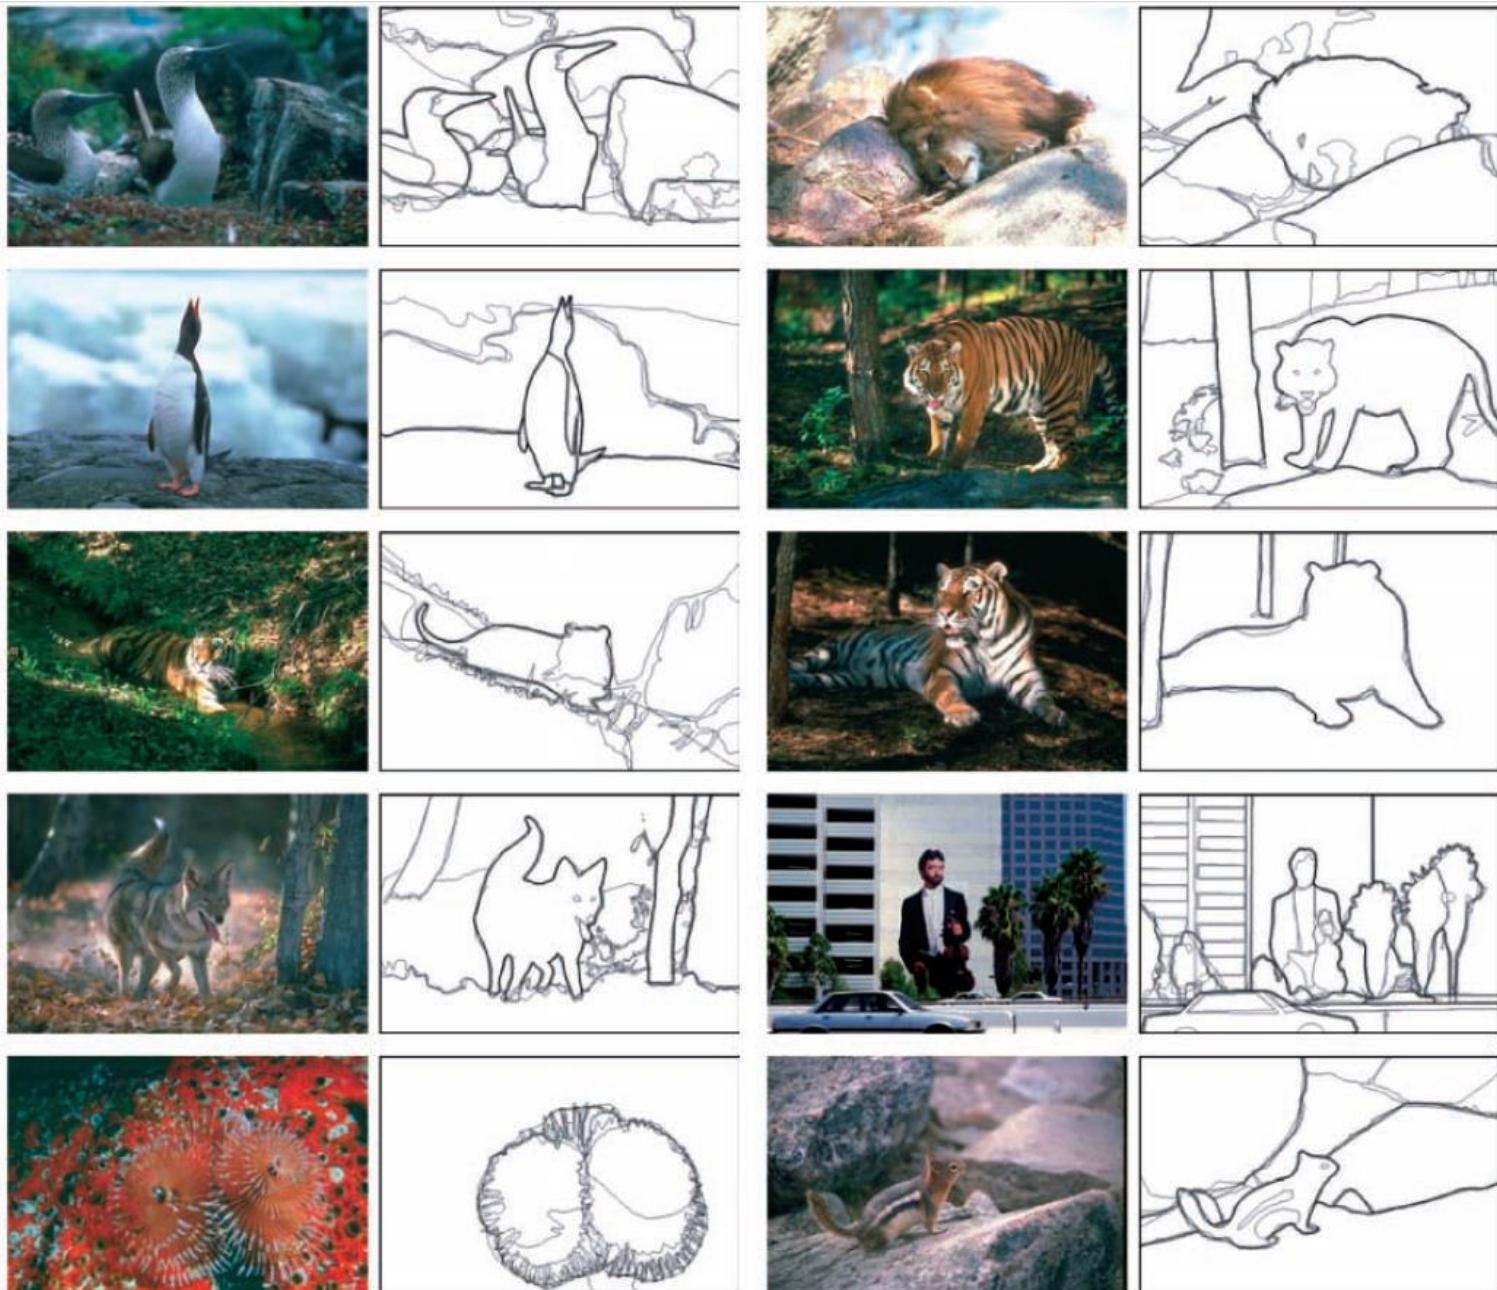
gradient magnitude



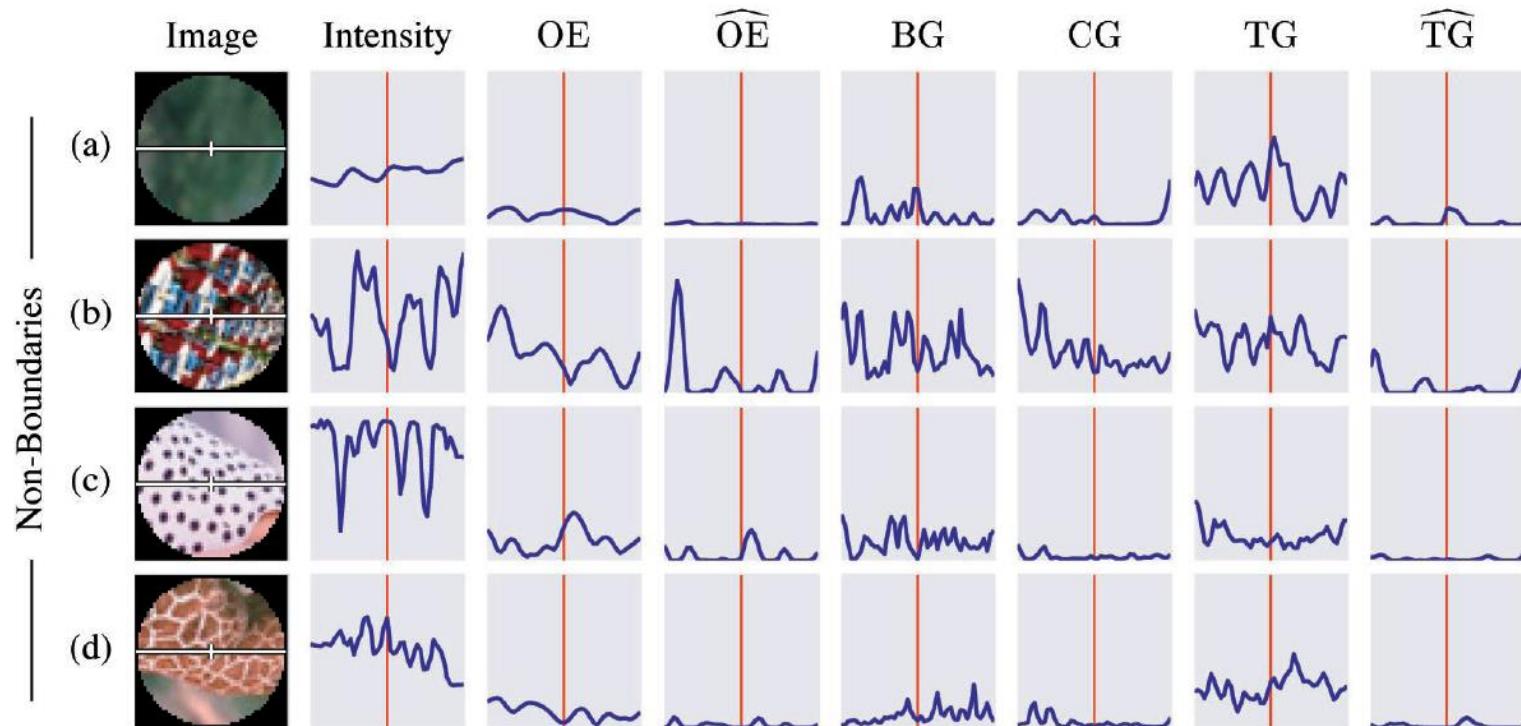
- Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

**Learn from  
humans which  
combination of  
features is  
most indicative  
of a “good”  
contour?**



# What features are responsible for perceived edges?

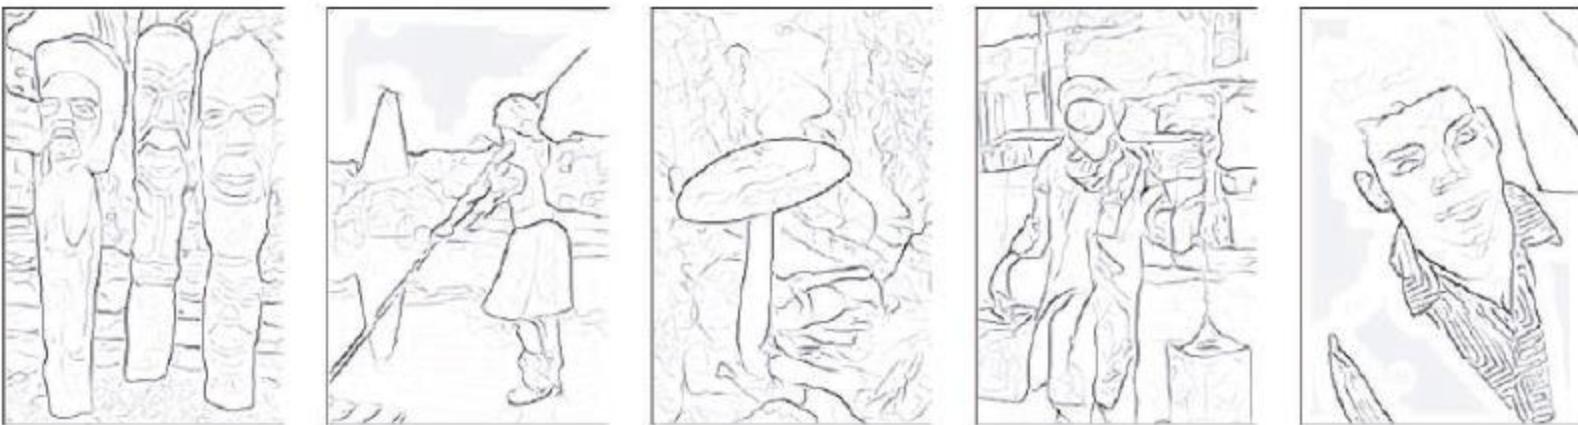


**Feature profiles (oriented energy, brightness, color, and texture gradients) along the patch's horizontal diameter**

Image



BG+CCG+TG



Human



# Summary

- Filters allow local image neighborhood to influence our description and features
  - Smoothing to reduce noise
  - Derivatives to locate contrast, gradient
- Convolution properties will influence the efficiency with which we can process images.
  - Associative
  - Filter separability
- Edge detection processes the image gradient to find curves, or chains of edgels.

# Slide Credits

- Slides Kristen Grauman, Steve Seitz
- and others, as marked...